



高性能计算与云计算

第九讲 CUDA编程

目录

➤ GPU微处理器简介

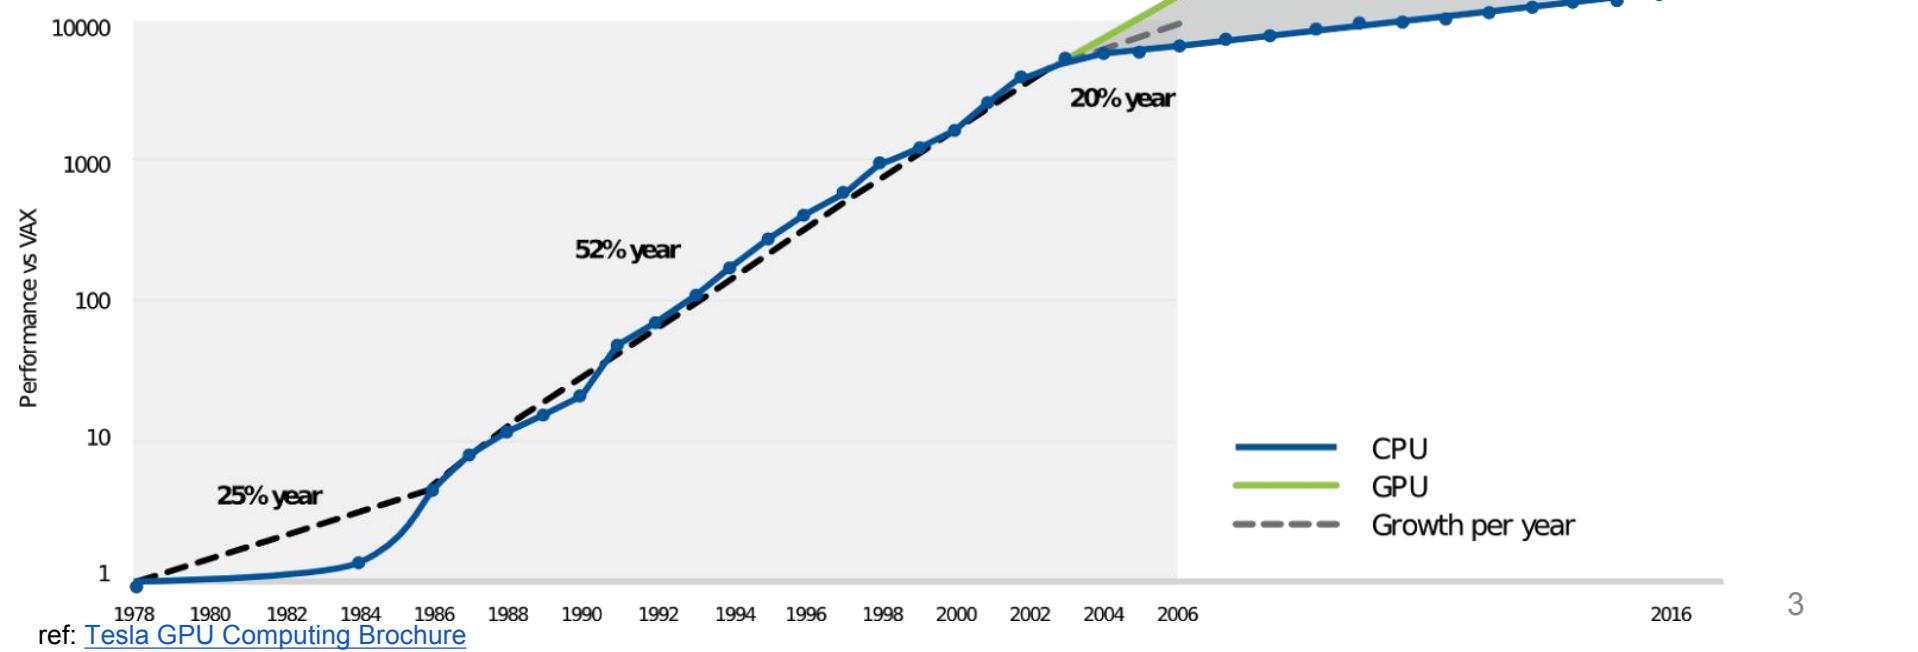
- ✓ GPU与CPU的比较
- ✓ GPU应用
- ✓ GPU举例

➤ CUDA编程

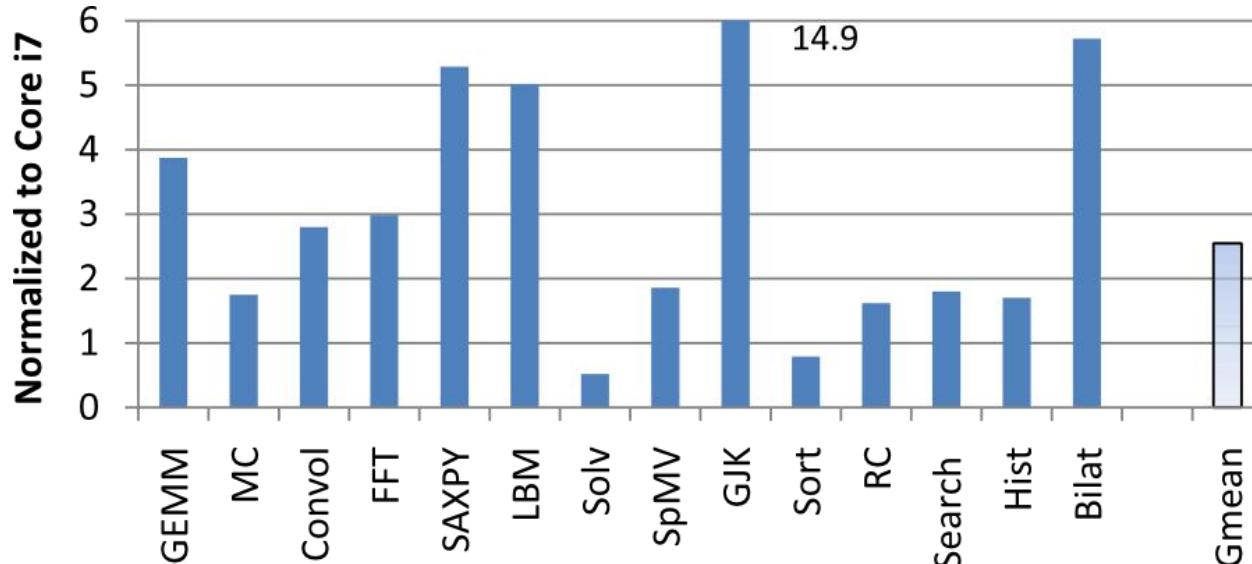
The GPU-CPU Gap (by NVIDIA)

Conventional CPU computing architecture can no longer support the growing HPC needs.

Source: Hennessy & Patterson, CAAQA, 4th Edition.

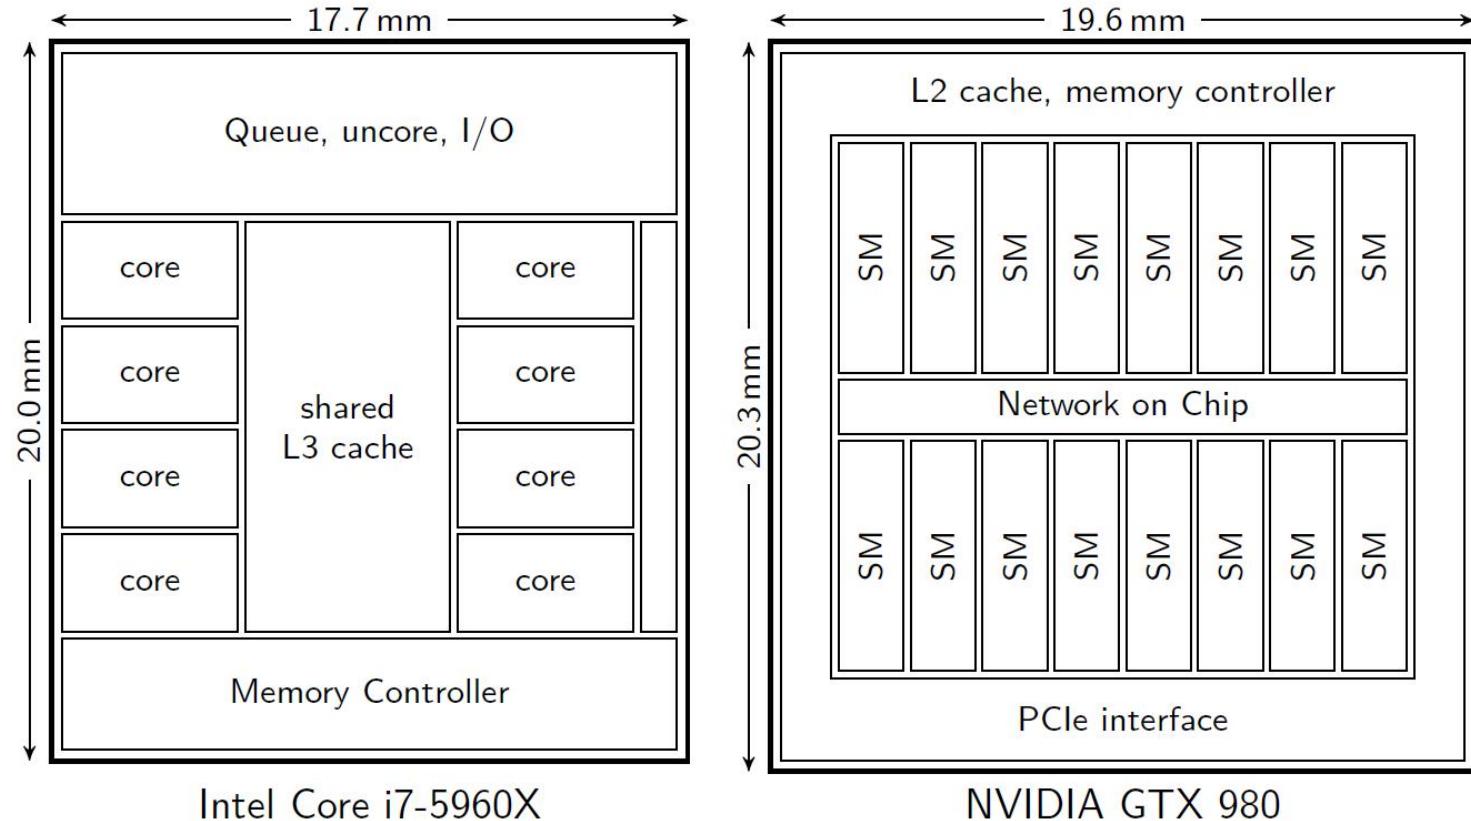


The GPU-CPU Gap (by Intel, 2010)



GTX280 v.s. Core i7 performance

CPU vs. GPU – 芯片面积



CPU与GPU的比较

- 传统CPU
 - ✓ 将芯片的大部分区域用于片上缓存；
 - ✓ CPU core比较重，用来处理非常复杂的控制逻辑，以优化串行程序执行。
- GPU
 - ✓ 将芯片上的大部分区域用于计算逻辑；
 - ✓ GPU Core比较轻，用于优化具有简单控制逻辑的数据并行任务，注重并行程序的吞吐量。
- 多核 (multi-core) vs. 众核 (many-core)
多核：CPU ; 众核：GPU

目录

➤ GPU微处理器简介

- ✓ GPU与CPU的比较
- ✓ GPU应用
- ✓ GPU举例

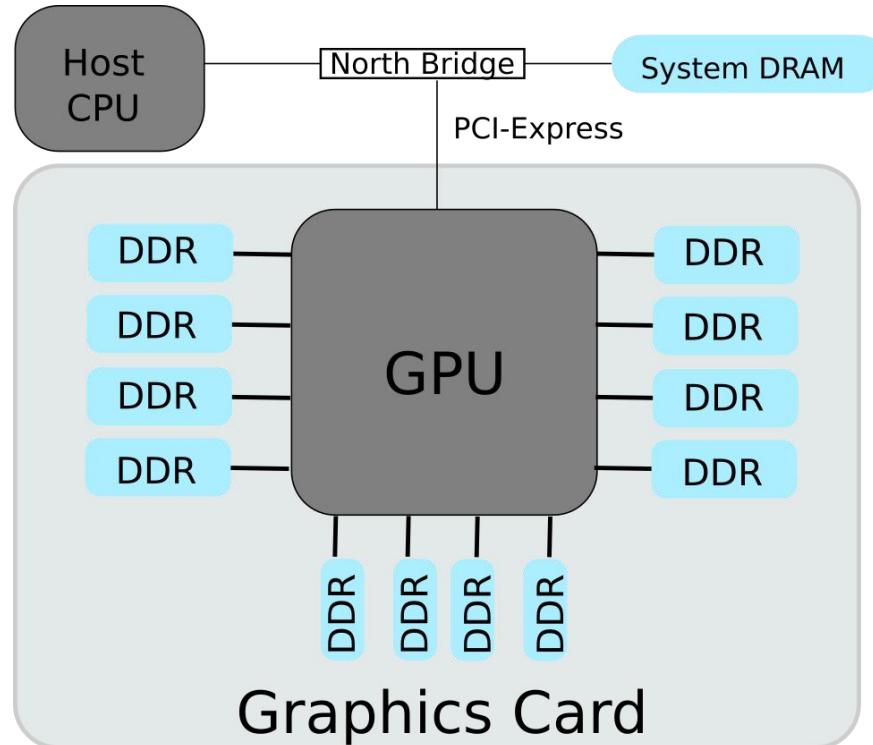
➤ CUDA编程

GPU的应用

- 历史上，GPU最初被设计用来专门处理并行图形计算问题，是**图形加速器**。
 - ✓ GPU使显卡减少了对CPU的依赖，并进行部分原本CPU的工作，尤其是在3D图形处理时
- 目前，GPU已演化成一个强大的、多用途的、完全可编程的，以及**任务和数据并行的处理器**。
- **异构架构**：GPU作为CPU的协处理器，GPU通过PCIe总线与CPU相连。

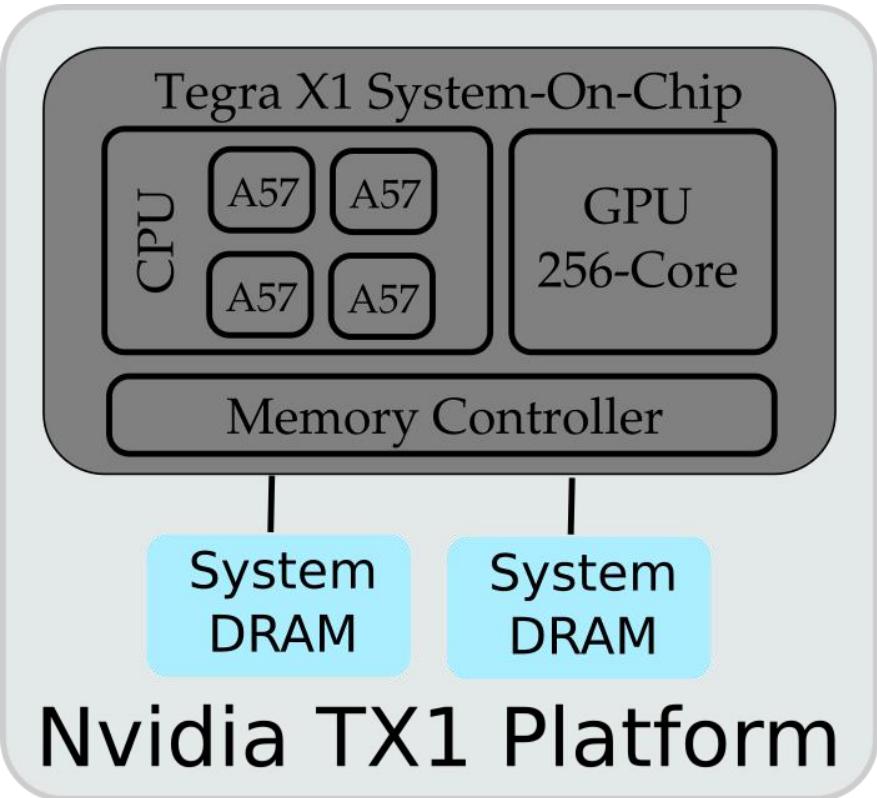
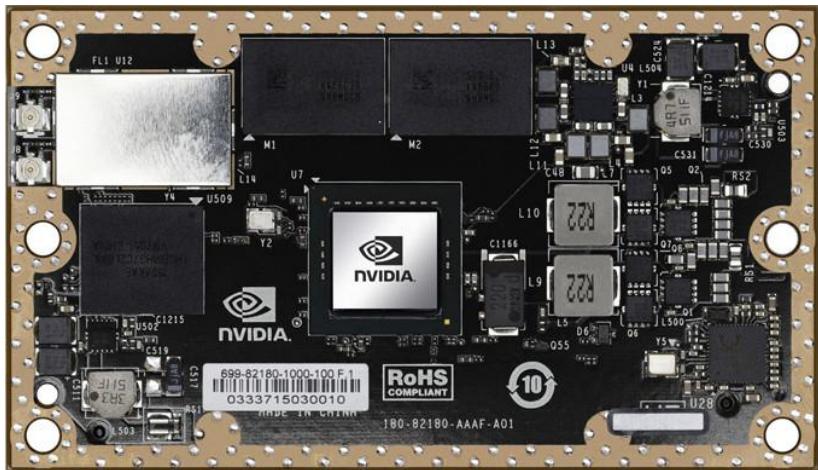
显卡中的GPU

Image: Nvidia GTX 980



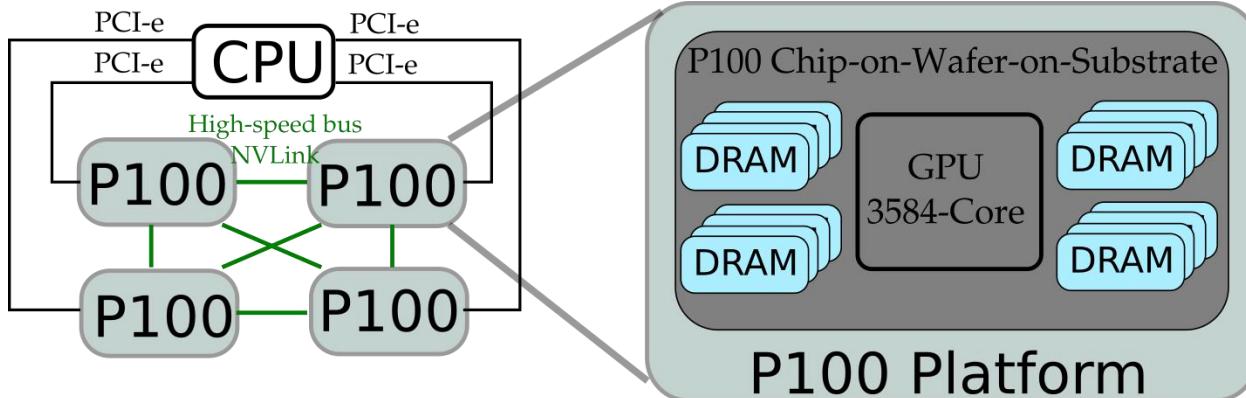
移动处理器中的GPU

Image: Nvidia Jetson TX1 (Tegra X1 SOC)

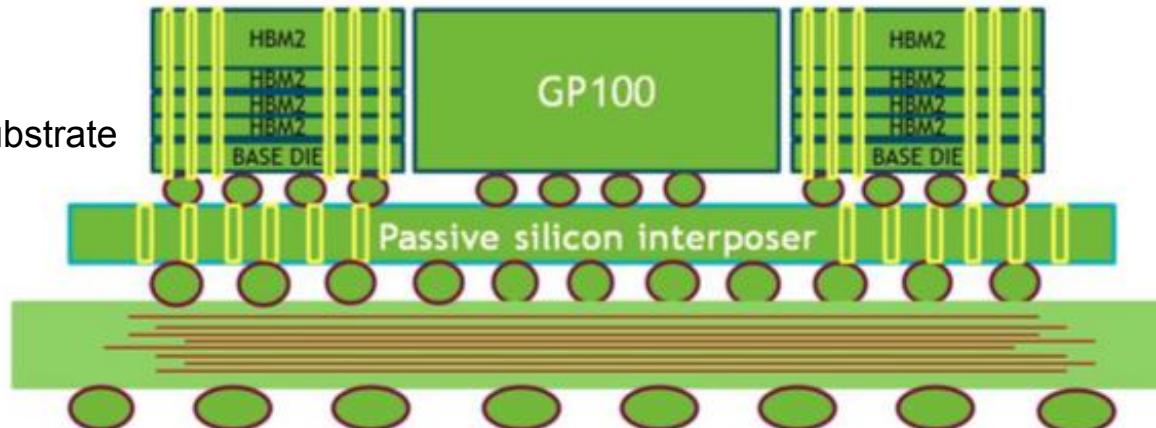


高性能计算机中的GPU

Image: Nvidia P100
(Pascal Architecture)



Chip-on-Wafer-on-Substrate



高性能计算机中的GPU

TOP500 supercomputer list in Jun. 2019. (<http://www.top500.org>)

Rank	System	Cores	Rmax (TFlop/s)	Rpeak (TFlop/s)	Power (kW)
1	Summit - IBM Power System AC922, IBM POWER9 22C 3.07GHz, NVIDIA Volta GV100 Dual-rail Mellanox EDR Infiniband , IBM DOE/SC/Oak Ridge National Laboratory United States	2,397,824	143,500.0	200,794.9	9,783
2	Sierra - IBM Power System S922LC, IBM POWER9 22C 3.1GHz, NVIDIA Volta GV100 Dual-rail Mellanox EDR Infiniband , IBM / NVIDIA / Mellanox DOE/NNSA/LLNL United States	1,572,480	94,640.0	125,712.0	7,438
3	Sunway TaihuLight - Sunway MPP, Sunway SW26010 260C 1.45GHz, Sunway , NRCPC National Supercomputing Center in Wuxi China	10,649,600	93,014.6	125,435.9	15,371

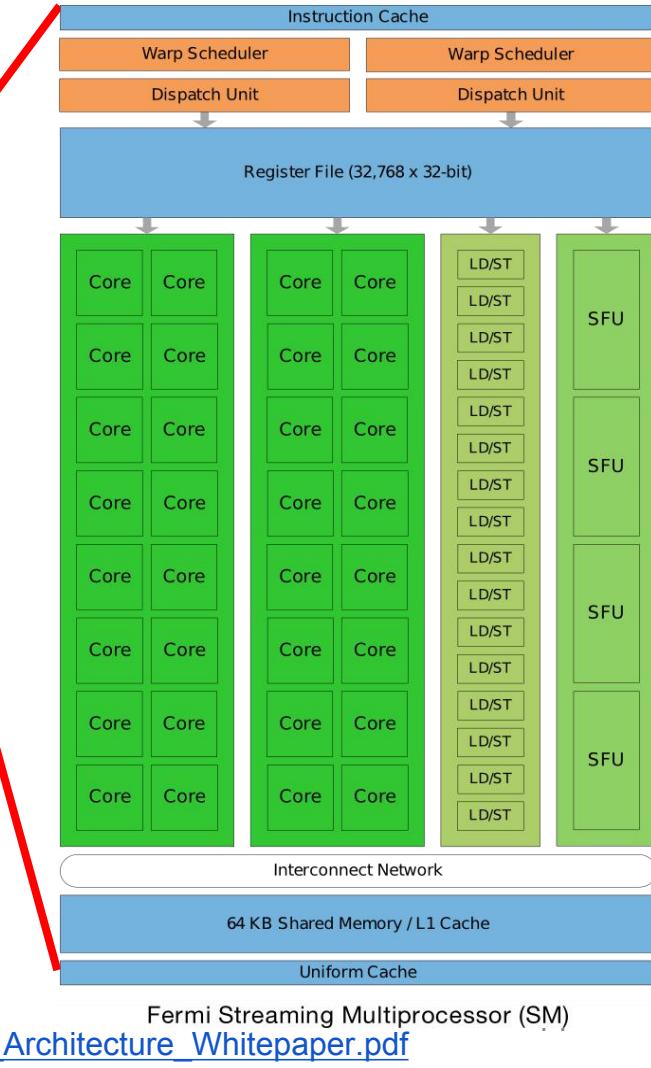
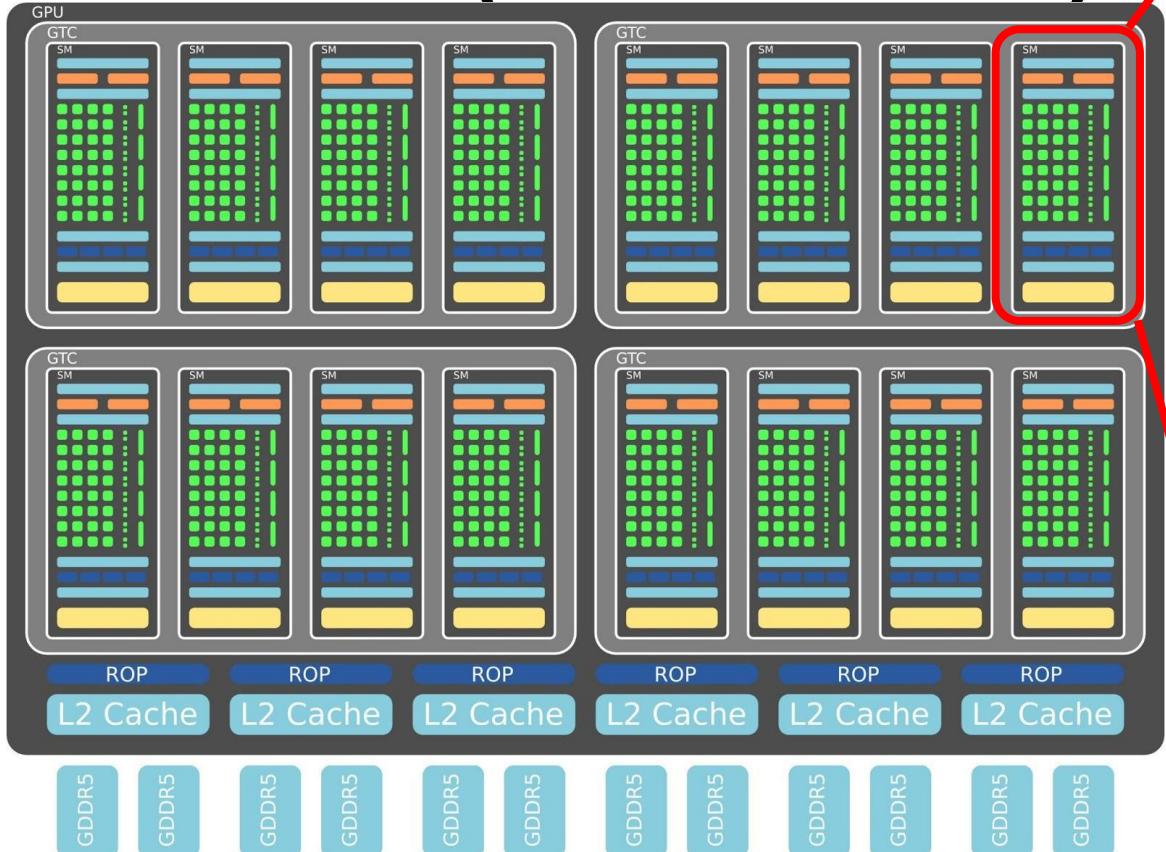
目录

➤ GPU微处理器简介

- ✓ GPU与CPU的比较
- ✓ GPU应用
- ✓ GPU举例

➤ CUDA编程

NVIDIA (Fermi 架构)



ref: http://www.nvidia.com/content/PDF/fermi_white_papers/NVIDIA_Fermi_Compute_Architecture_Whitepaper.pdf

晶体管数目

ref: http://en.wikipedia.org/wiki/Transistor_count

SOC

处理器	晶体管	工艺技术
61-Core Xeon Phi	5,000,000,000	22nm
22-core Xeon Broadwell-E5	7,200,000,000	14nm
32-Core Sparc M7	10,000,000,000+	20nm

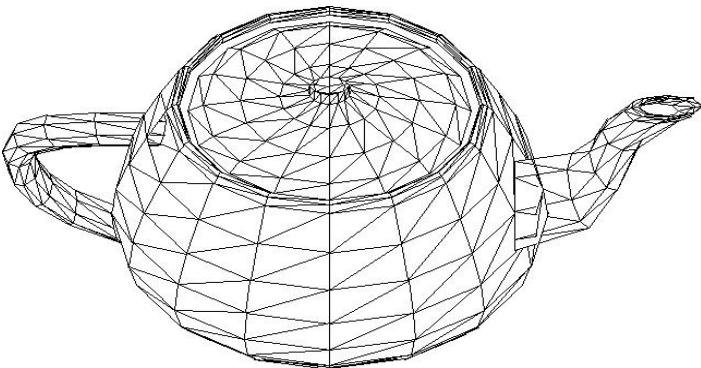
GPU

处理器	晶体管	工艺技术
Nvidia GP100 Pascal	15,300,000,000	16nm

FPGA

FPGA	晶体管	工艺技术
Virtex-Ultrascale XCVU440	20,000,000,000+	20nm
Stratix 10 10GX5500	30,000,000,000+	14nm

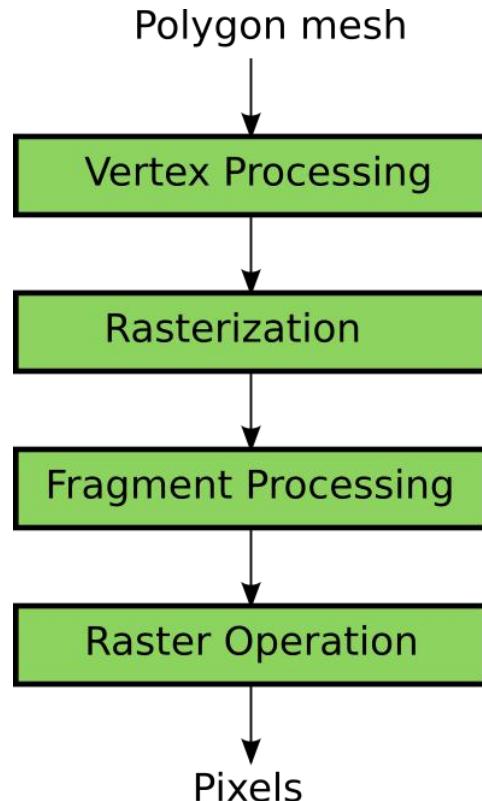
举例：153亿个晶体管的作用



Render triangles.
Billions of triangles per second.

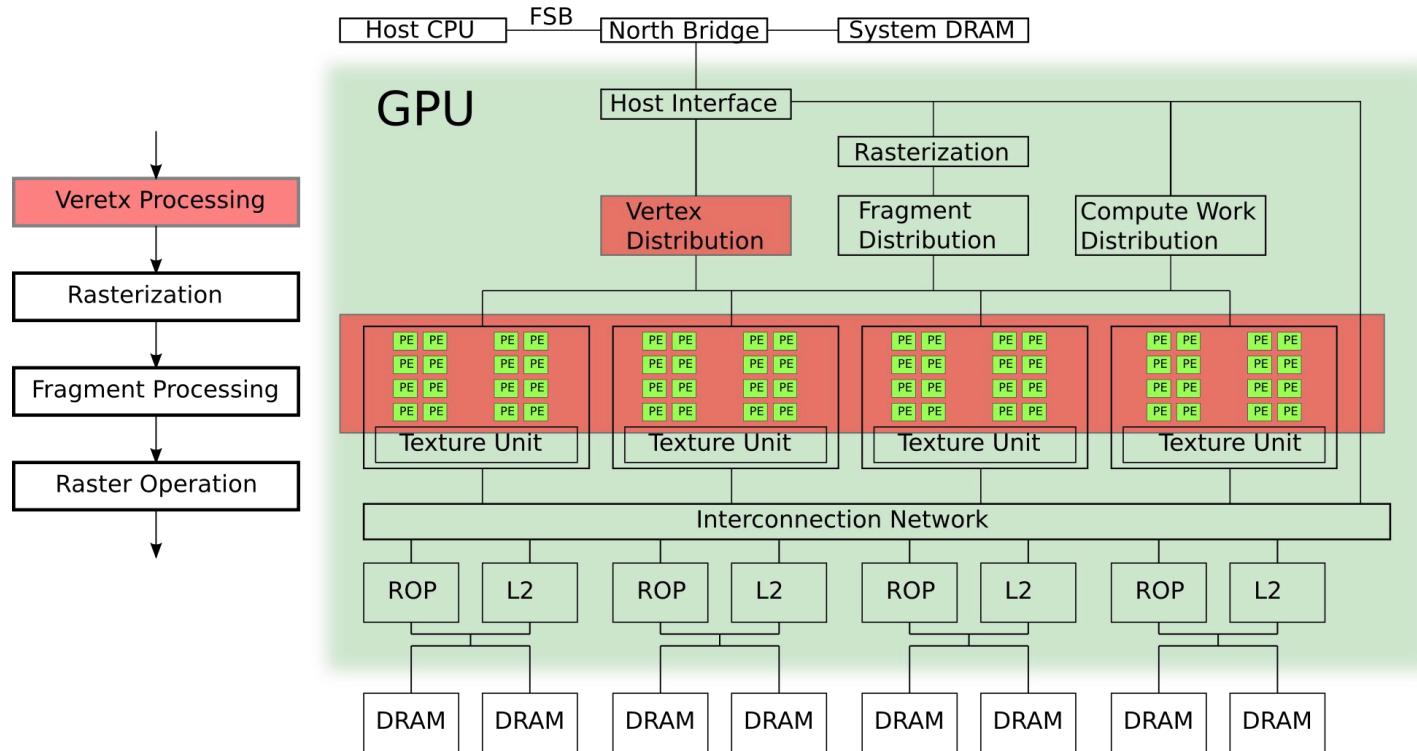


图形流水线 (The Graphics Pipeline)

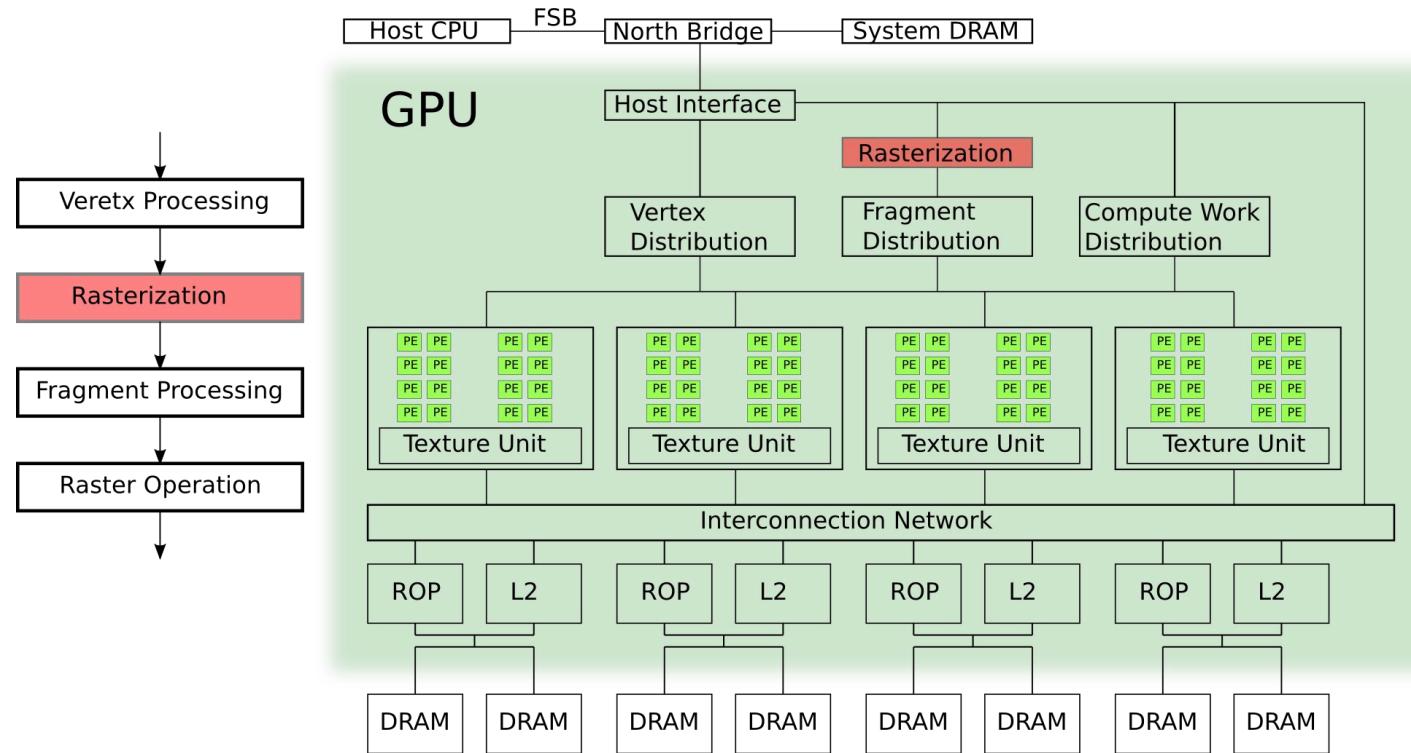


- ◆ 图形流水线将3D模型表示为多边形网格，并根据当前帧的视图渲染模型。
- ◆ 基本步骤：定点处理，光栅化，片元处理，光栅操作
- ◆ 现代图形流水线更加复杂。

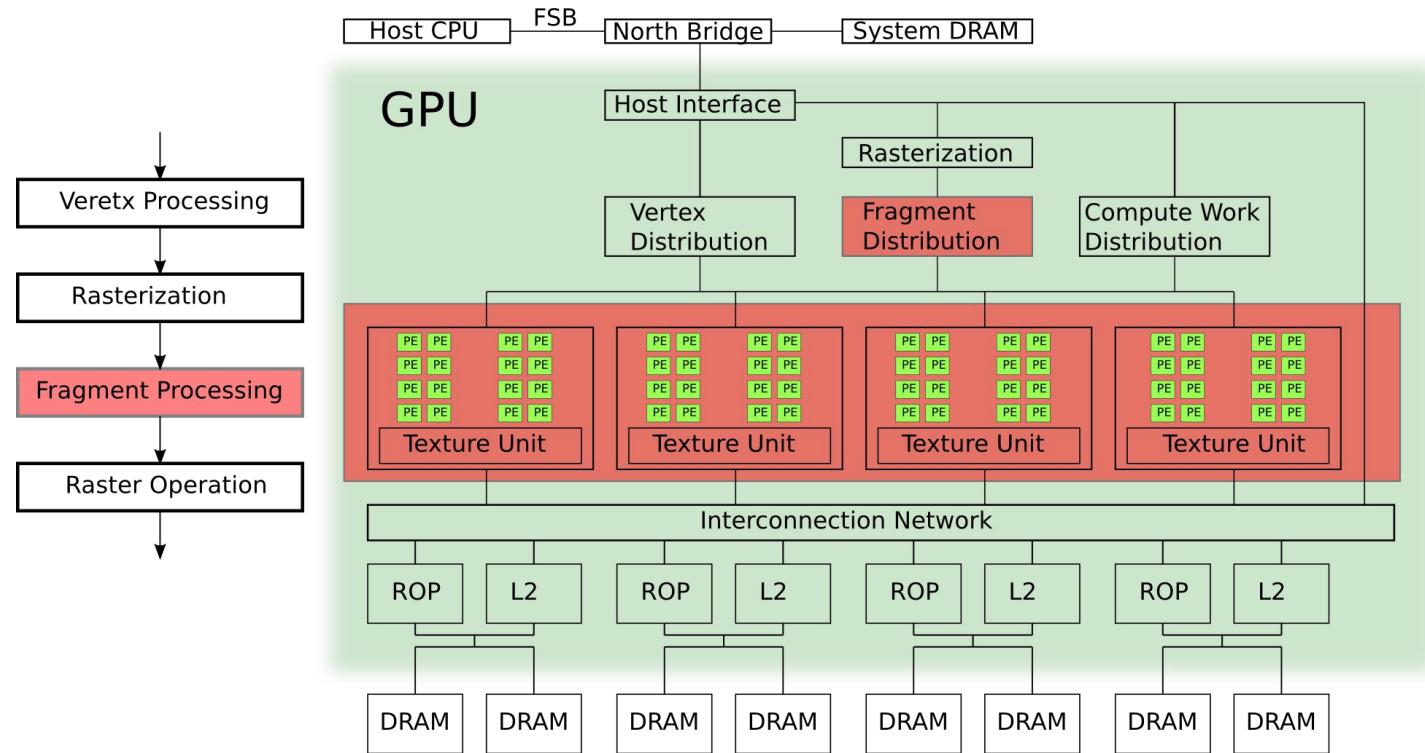
GPU的图形流水线 (1)



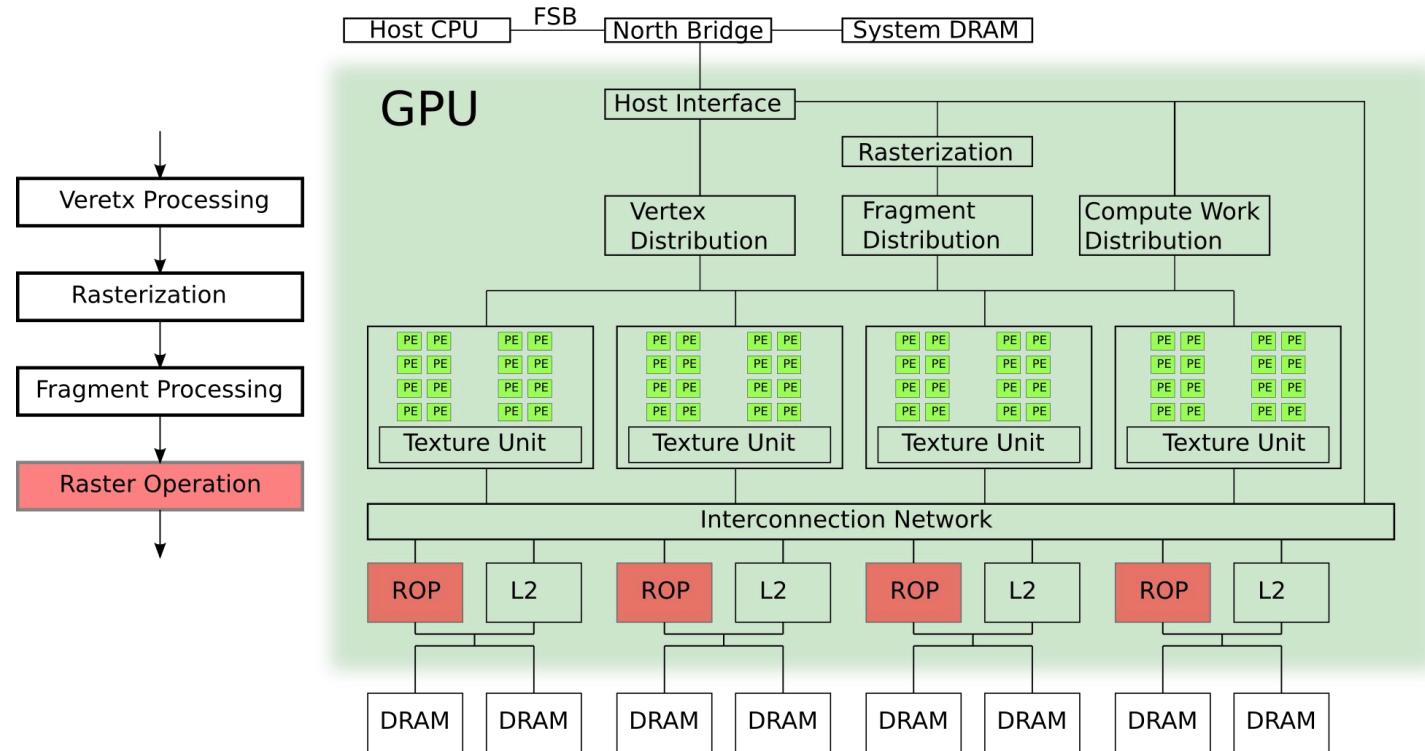
GPU的图形流水线 (2)



GPU的图形流水线 (3)



GPU的图形流水线 (4)



GPUs内部区域

GPUs are designed to match the workload of 3D graphics.

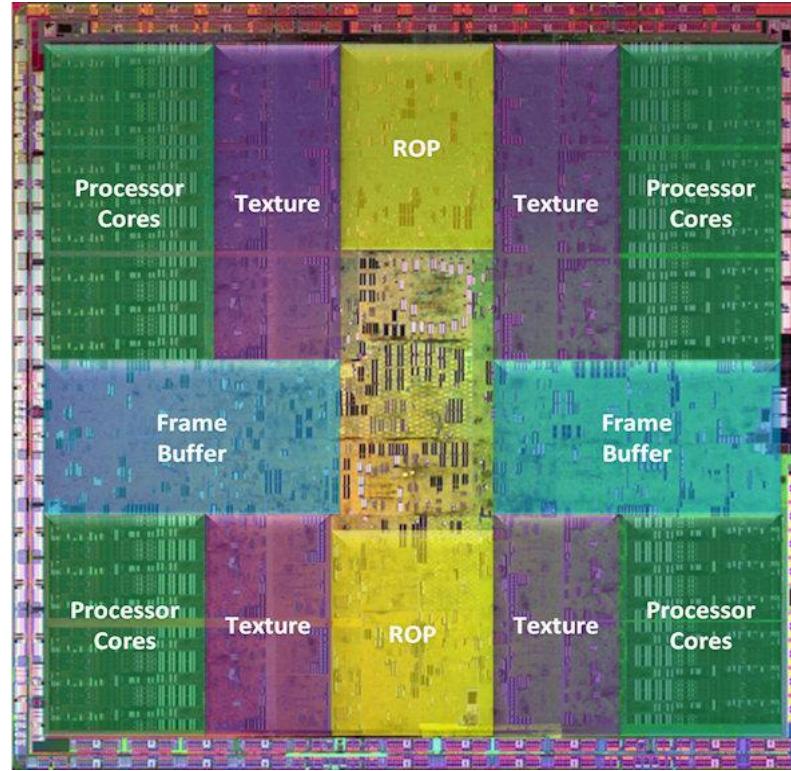
Die photo of GeForce GTX 280 (source: NVIDIA)

Texture:

For fragment processing.

ROP & Frame Buffer:

For raster operation.



J. Roca, et al. "Workload Characterization of 3D Games", IISWC 2006, [link](#)

T. Mitra, et al. "Dynamic 3D Graphics Workload Characterization and the Architectural Implications", Micro 1999, [link](#)

目录

- GPU微处理器简介
- CUDA编程
 - 引言：开发平台、异构计算
 - CUDA编程模型：线程、线程块、线程网格
 - CUDA执行模型：流多处理器和线程束
 - 内存层次结构
 - 编程实例

GPU计算的历史

- ▶ 2001/2002 - 研究人员把GPU当做数据并行协处理器
 - ▶ *GPGPU* 这个新领域从此诞生
- ▶ 2007 - NVIDIA 发布 *CUDA*
 - ▶ *CUDA* - 全称Compute Uniform Device Architecture
统一计算设备架构
 - ▶ GPGPU 发展成 *GPU Computing*
- ▶ 2008 - Khronos 发布 *OpenCL* 规范



GPU主流开发平台

➤ CUDA：

Compute Unified Device Architecture，2006年英伟达发布，将GPU计算由传统图形处理领域扩展到通用计算领域，支持c, c++, Fortran和Python等语言，提供2组API（一个底层，一个更高层），是实现更高层第三方API和库的基础。CUDA只针对英伟达硬件。

➤ OpenCL

Open computing Language开放计算语言，可以在如GPU、CPU、DSP和其他处理器等众多计算平台上进行编程的开放标准，英伟达和AMD都支持OpenCL。与CUDA编程模型相似。

GPU主流开发平台（2）

➤ OpenACC

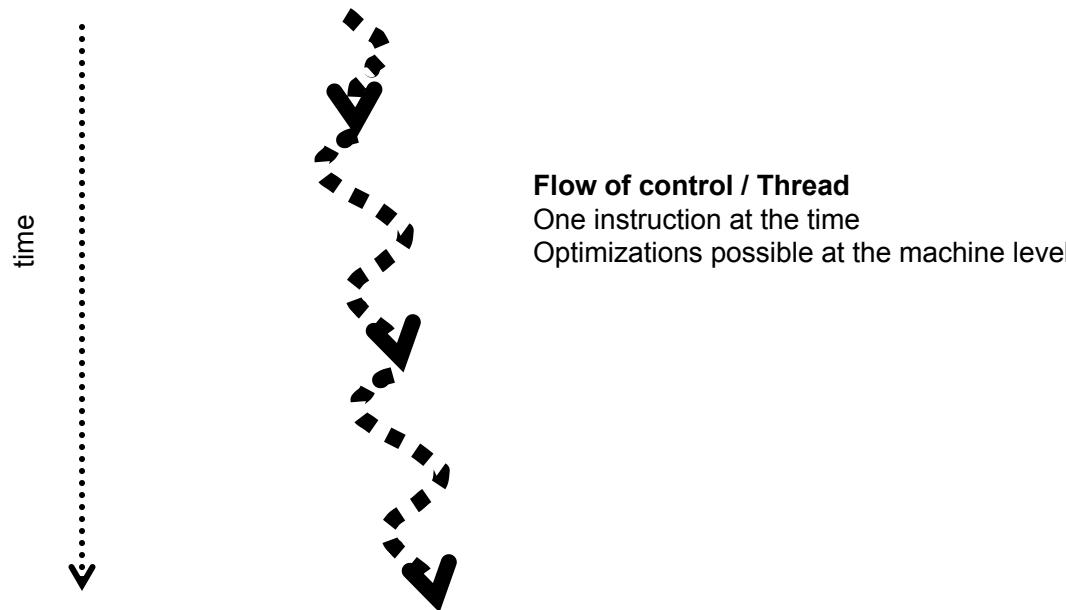
是一个API的开放规范，允许使用编译制导语句（如`#pragma acc`，风格与OpenMP类似），根据用户提示自动将计算映射到GPU或者多核芯片上。

➤ C++AMP

是微软基于DirectX 11开发的技术，允许C++代码根据程序员提供的大量指令或者语言扩展在CPU或者GPU计算平台上透明执行。

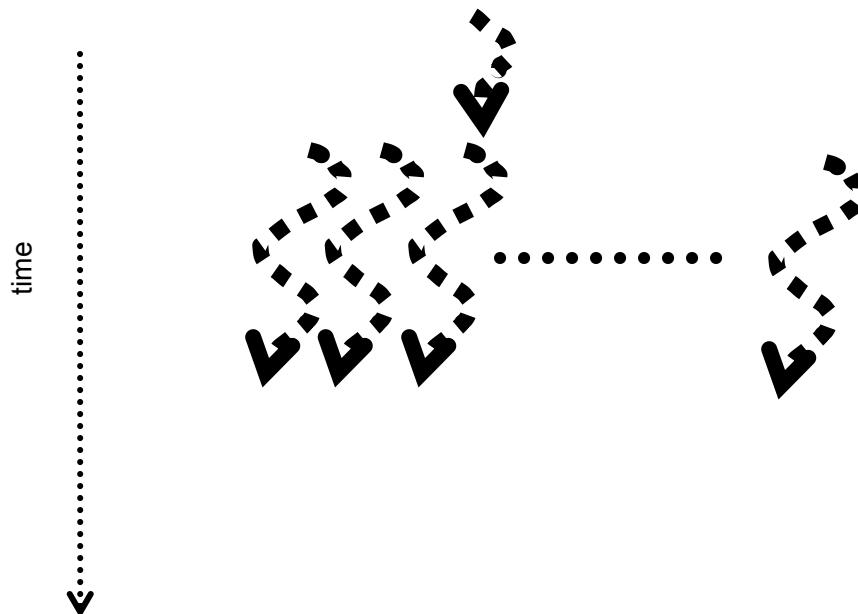
- 回顾: Sequential Execution Model

```
int a[N]; // N is large  
for (i = 0; i < N; i++)  
    a[i] = a[i] * fade;
```



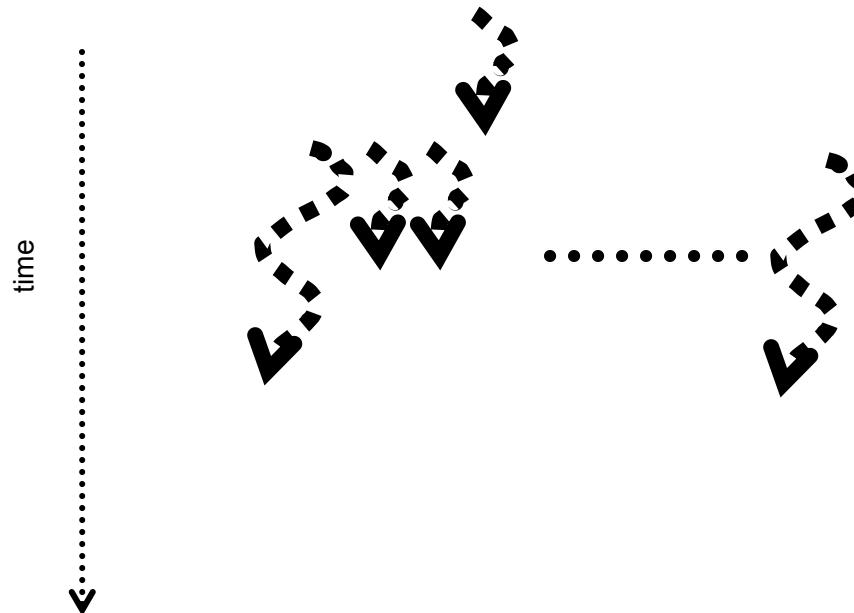
- 回顾: Data Parallel Execution Model / SIMD

```
int a[N]; // N is large  
for all elements do in parallel  
    a[index] = a[index] * fade;
```



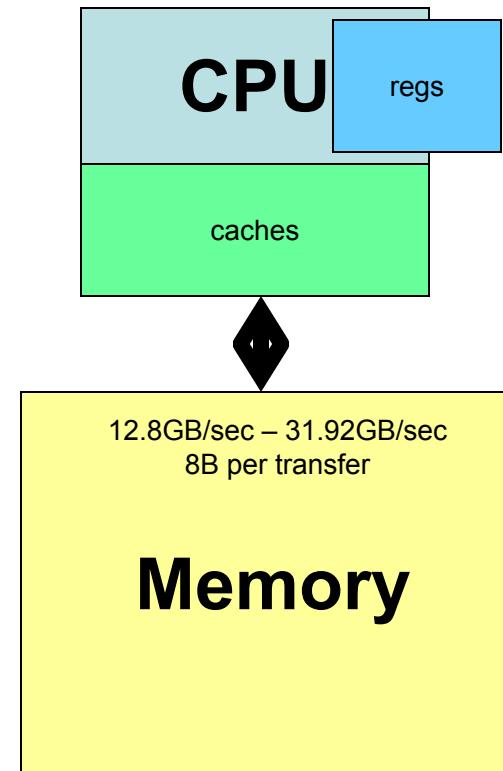
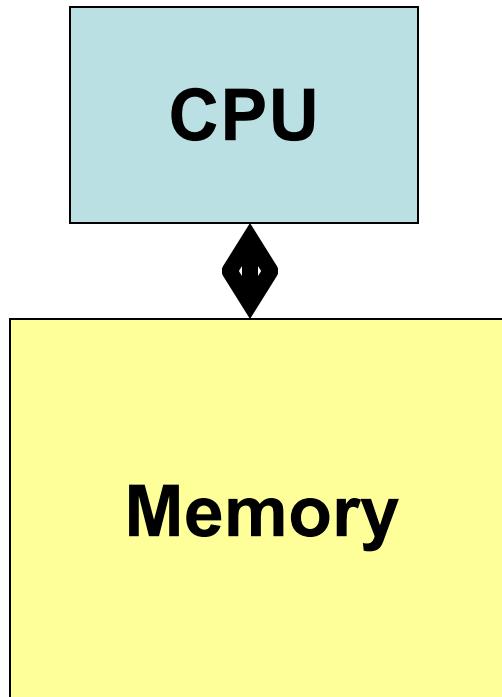
- 回顾: **Single Program Multiple Data / SPMD**

```
int a[N]; // N is large  
for all elements do in parallel  
if (a[i] > threshold) a[i] *= fade;
```



- 回顾：程序员视角 – 典型系统

如特别关注性能



异构计算

- 术语:
 - 主机(Host) CPU和它的内存 (主机内存)
 - 设备(Device) GPU和它的内存 (设备内存)



Host

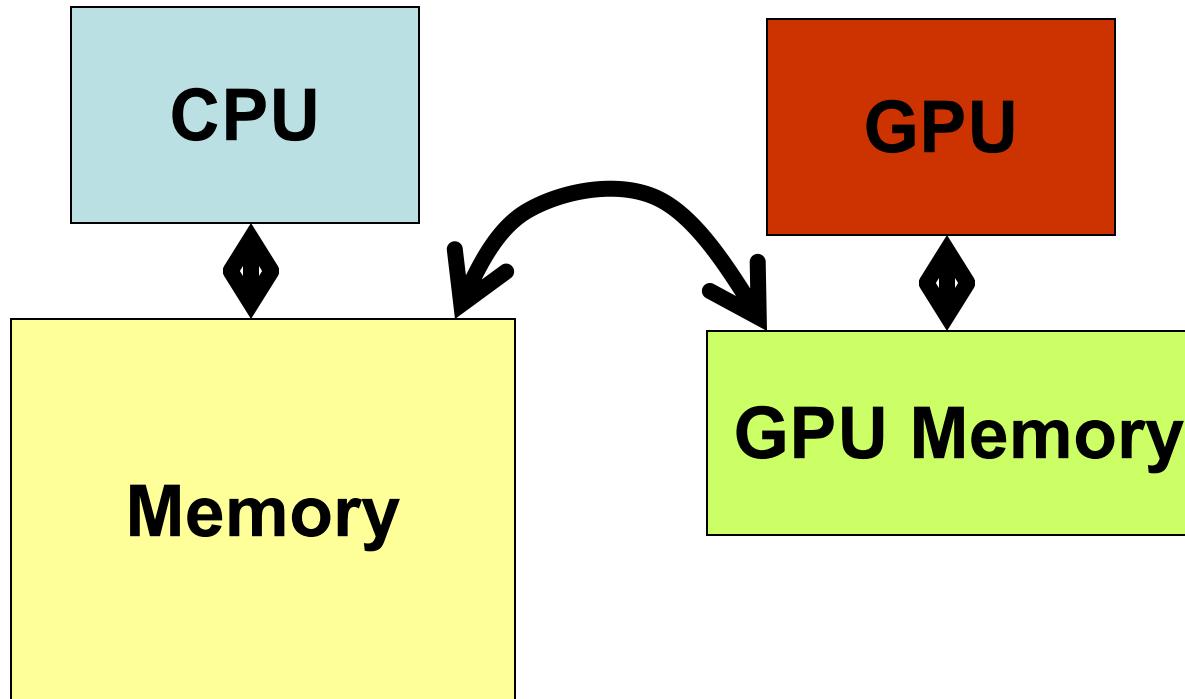


Device

- 异构计算 (2)

- 主机 (Host)

- 设备 (Device)



异构计算：全局串行局部并行

```
#include <iostream>
#include <algorithm>

using namespace std;

#define N 1024
#define RADIUS 3
#define BLOCK_SIZE 16

__global void stencil_1d(int *in, int *out) {
    __shared__ int temp[BLOCK_SIZE + 2 * RADIUS];
    int index = threadIdx.x * blockDim.x;
    int tempIndex = index - RADIUS;
    if (tempIndex < 0) tempIndex = 0;
    if (tempIndex >= RADIUS) tempIndex = RADIUS;
    if (tempIndex <= index) {
        temp[tempIndex - RADIUS] = in[index - RADIUS];
        temp[tempIndex + RADIUS] = in[index + RADIUS];
    }
    __syncthreads();
    if (index <= RADIUS) return;
    int result = 0;
    for (int offset = -RADIUS; offset <= RADIUS; offset++)
        result += temp[index + offset];
    if (tempIndex <= index) {
        out[index] = result;
        return;
    }
    __syncthreads();
}

void fill(int *x, int n) {
    for (int i = 0; i < n; i++)
        x[i] = i;
}

int main(void) {
    int *in, *out; // host copies of a, b, c
    int *d_in, *d_out; // device copies of a, b, c
    int size = (N + 2 * RADIUS) * sizeof(int);

    // Alloc space for host copies and setup values
    in = (int *)malloc(size);
    d_in = (int *)malloc(size);
    out = (int *)malloc(size);
    fill(in, N + 2 * RADIUS);

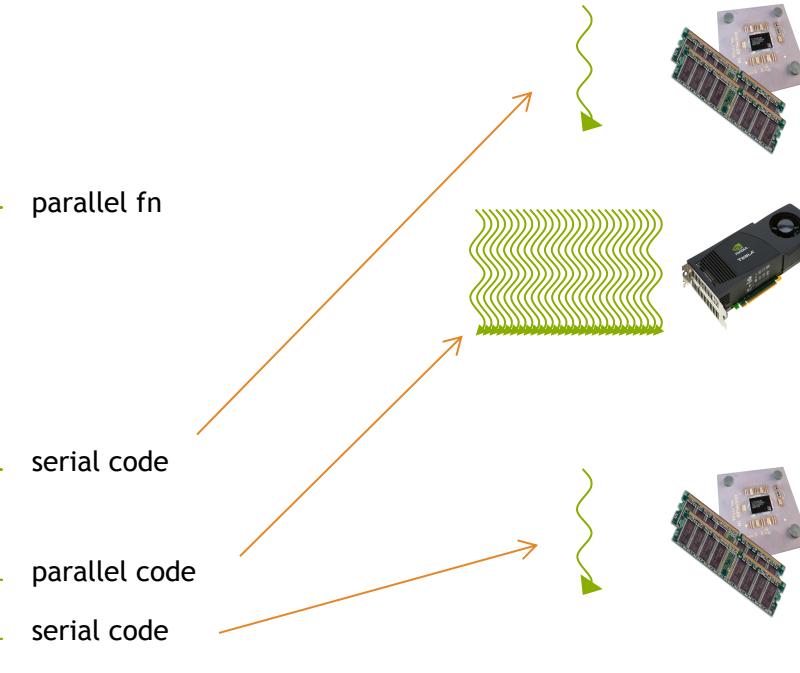
    // Alloc space for device copies
    cudaMalloc((void **) &d_in, size);
    cudaMalloc((void **) &d_out, size);

    // Copy to device
    cudaMemcpy(d_in, in, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_out, out, size, cudaMemcpyHostToDevice);

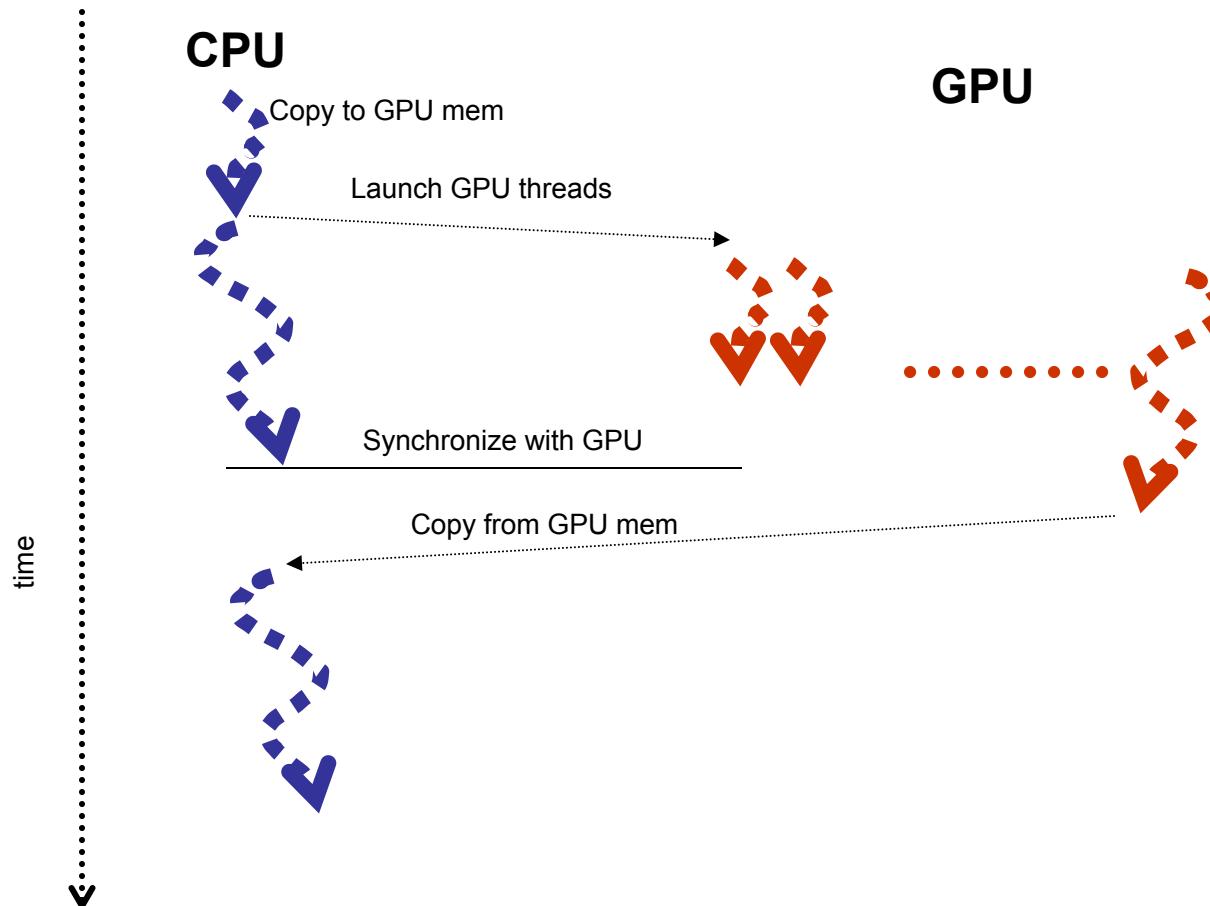
    // Launch stencil_1d() kernel on GPU
    stencil_1d<<N<<BLOCK_SIZE<<BLOCK_SIZE>>>(d_in + RADIUS, d_out + RADIUS);

    // Copy result back to host
    cudaMemcpy(out, d_out, size, cudaMemcpyDeviceToHost);

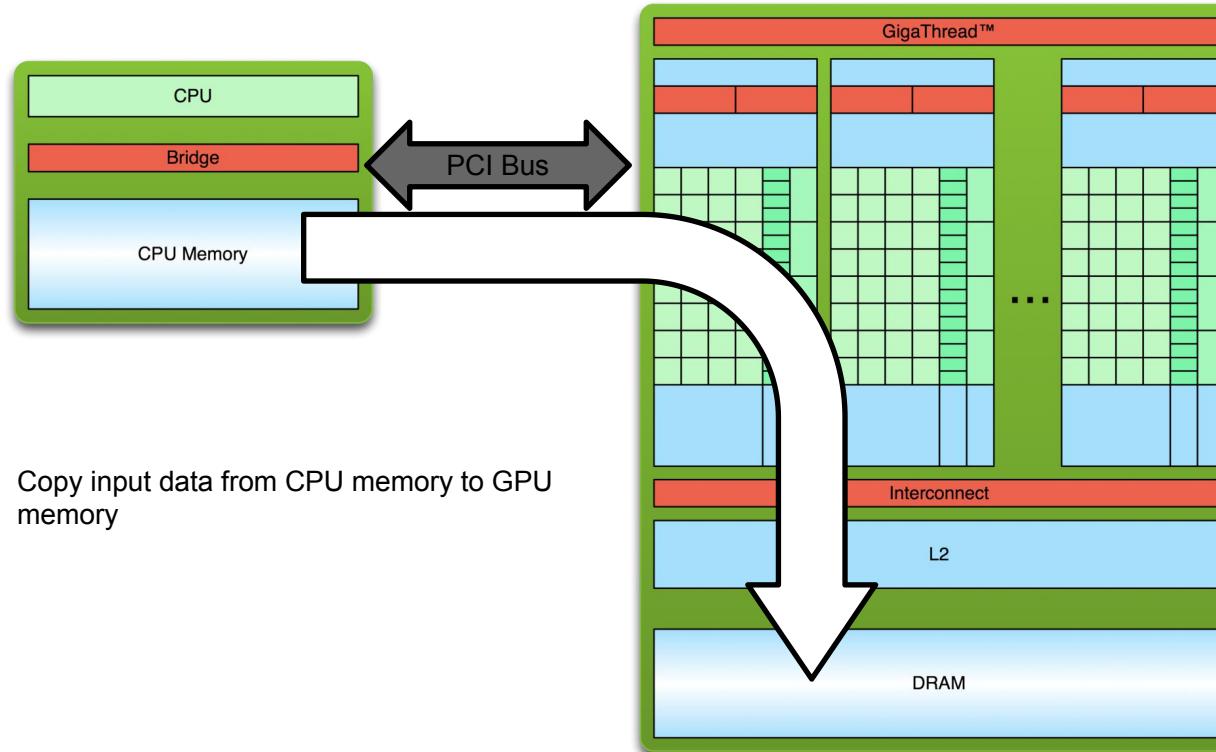
    // Cleanup
    free(in);
    free(d_in);
    free(d_out);
    return 0;
}
```



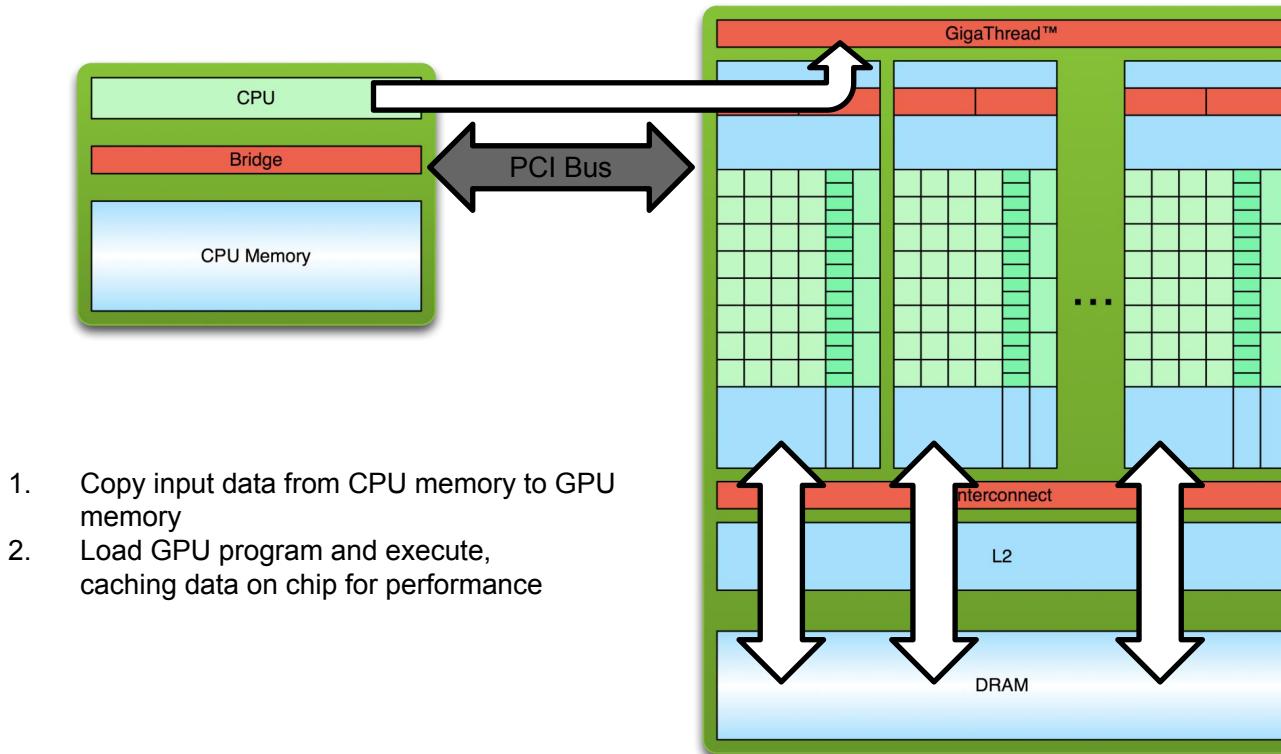
- 异构计算：全局串行局部并行（2）



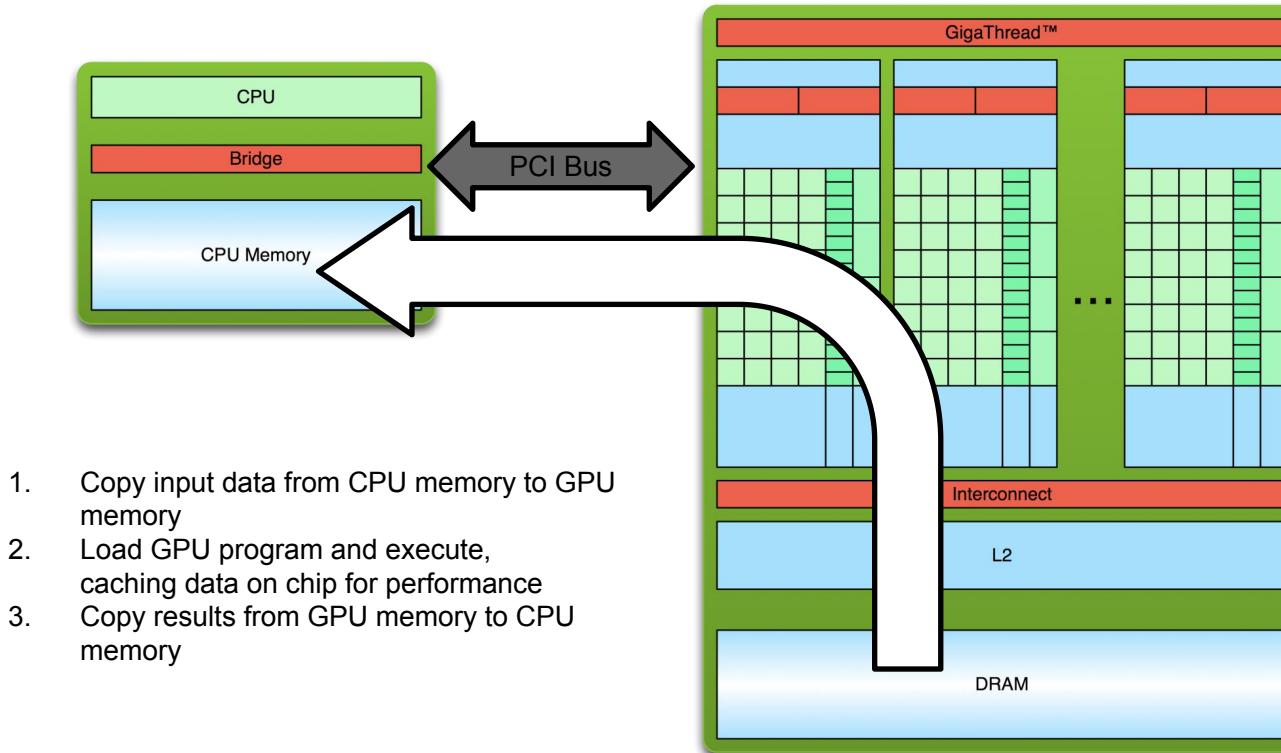
简单处理流程



简单处理流程 (2)



简单处理流程 (3)



目录

- GPU微处理器简介
- CUDA编程
 - 引言：开发平台、异构计算
 - **CUDA编程模型：线程、线程块、线程网格**
 - CUDA执行模型：流多处理器和线程束
 - 内存层次结构
 - 编程实例

Hello World

```
#include <stdio.h>
__global__ void helloFromGPU (void)
{
    printf("Hello World from GPU!\n");
}

int main(void)
{
// hello from cpu
printf("Hello World from CPU!\n");

    helloFromGPU <<<1, 10>>>();

    cudaDeviceReset();

    return 0;
}
```

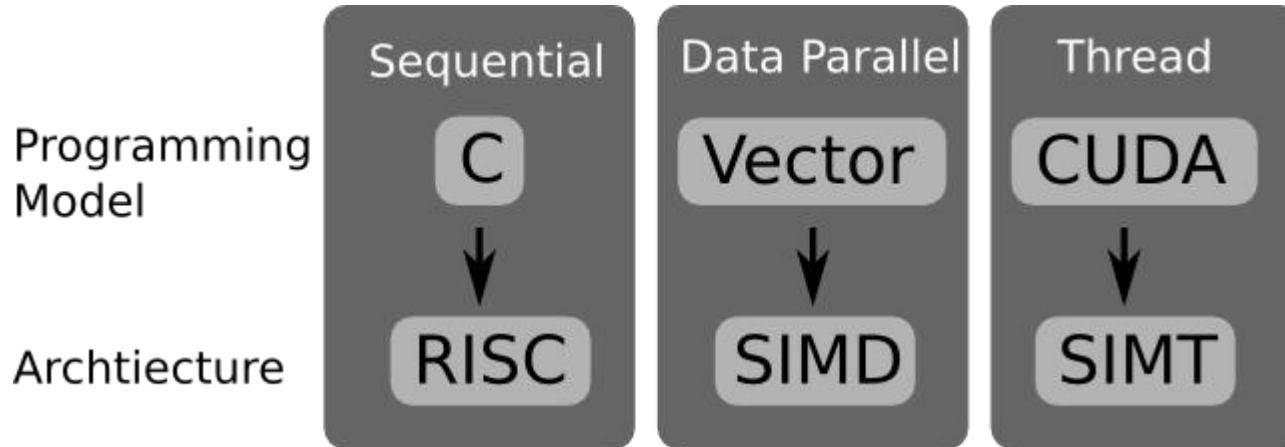
```
$nvcc –arch sm_20 hello.cu –o hello
$ ./hello
```

```
Hello World from CPU!
Hello World from GPU!
```

程序结构的5个步骤：

1. 分配GPU空间；
2. 将数据从CPU端复制到GPU端；
3. 调用CUDA kernel来执行计算；
4. 计算完成后将数据从GPU拷贝回CPU；
5. 清理GPU内存空间。

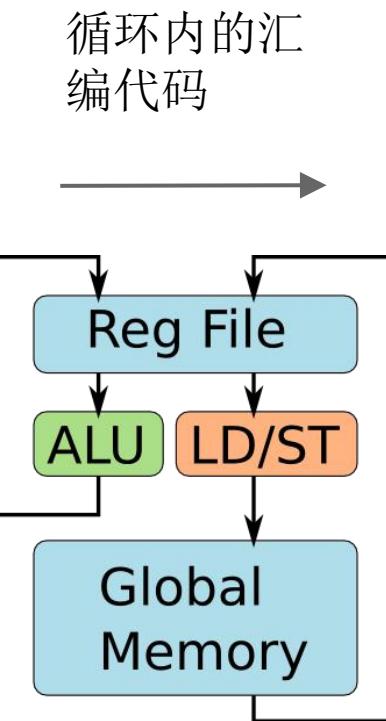
举例：GPU编程—编程模型与架构



回顾：C语言和RISC架构

回顾：C 语言和 RISC 架构

```
int A[2][4];
for(i=0;i<2;i++){
    for(j=0;j<4;j++){
        A[i][j]++;
    }
}
```

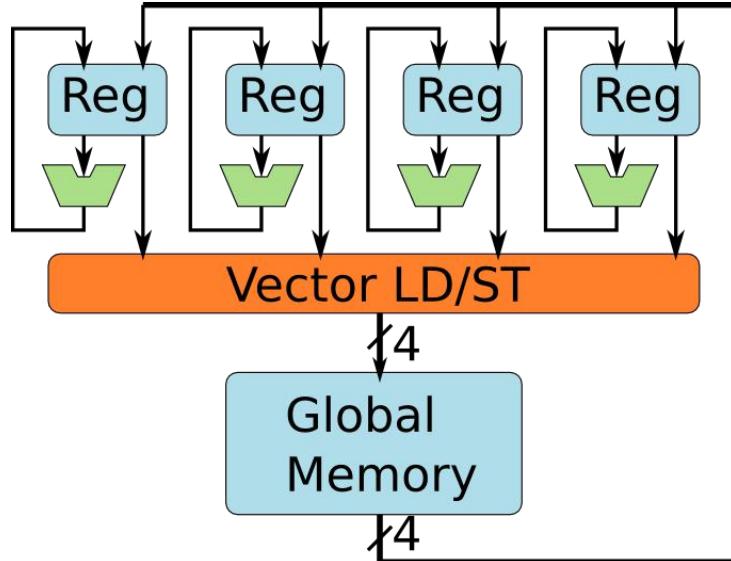


lw r0, 4(r1)
addi r0, r0, 1
sw r0, 4(r1)

Programmer's view of RISC

回顾：C 语言和 RISC 架构（2）

许多 CPUs 具有 **vector**（矢量）
单指令多数据流（SIMD）单元



- Programmer's view of a vector SIMD
- 如 Intel Streaming SIMD Extension(SSE)

回顾：C 语言和 RISC 架构（3）

矢量 SIMD 编程

内部循环展开成向量 SIMD：

```
int A[2][4];
for(i=0;i<2;i++){
    for(j=0;j<4;j++){
        A[i][j]++;
    }
}
```

```
int A[2][4];
for(i=0;i<2;i++){
    movups xmm0, [ &A[i][0] ] // load
    addps  xmm0,    xmm1      // add 1
    movups [ &A[i][0] ], xmm0   // store
}
```

与前面的例子一样，但每个 SSE 指令在 4 ALUs 上执行。

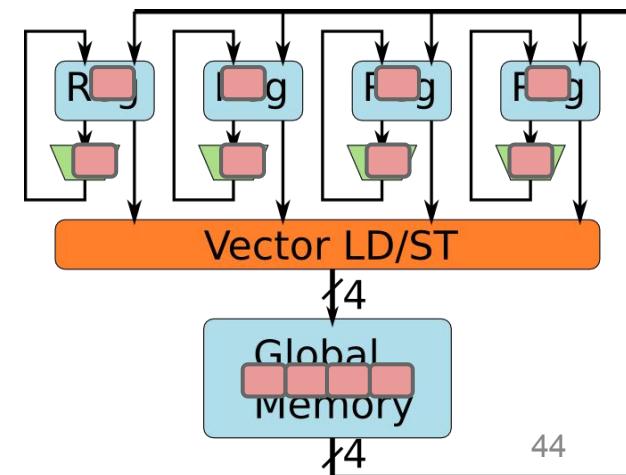
回顾：C 语言和 RISC 架构（4）

矢量程序的执行

```
int A[2][4];
for(i=0;i<2;i++){
```

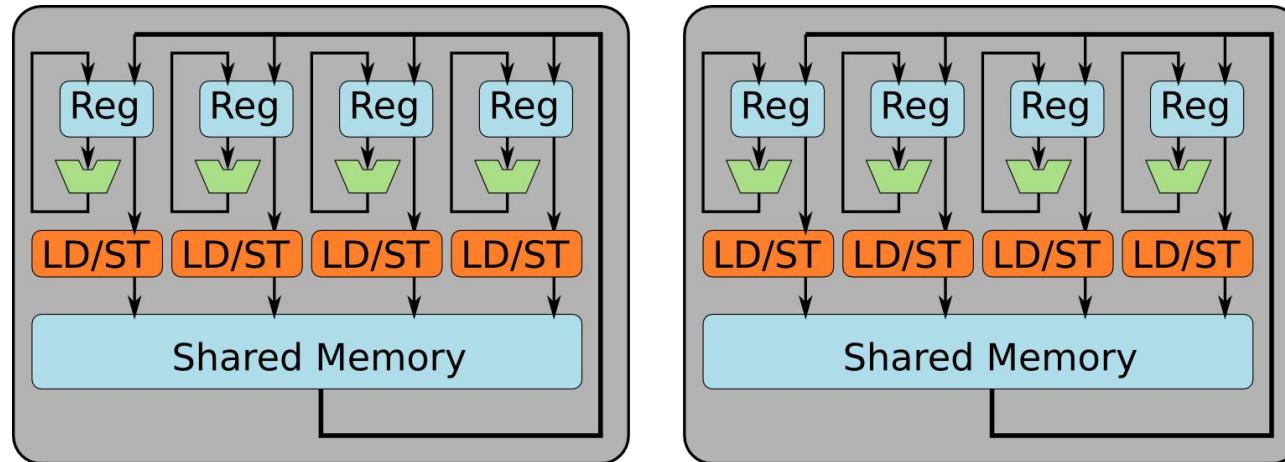
```
    movups xmm0,      [ &A[i][0] ] // load
    addps   xmm0,      xmm1       // add 1
    movups [ &A[i][0] ], xmm0    // store
}
```

```
}
```



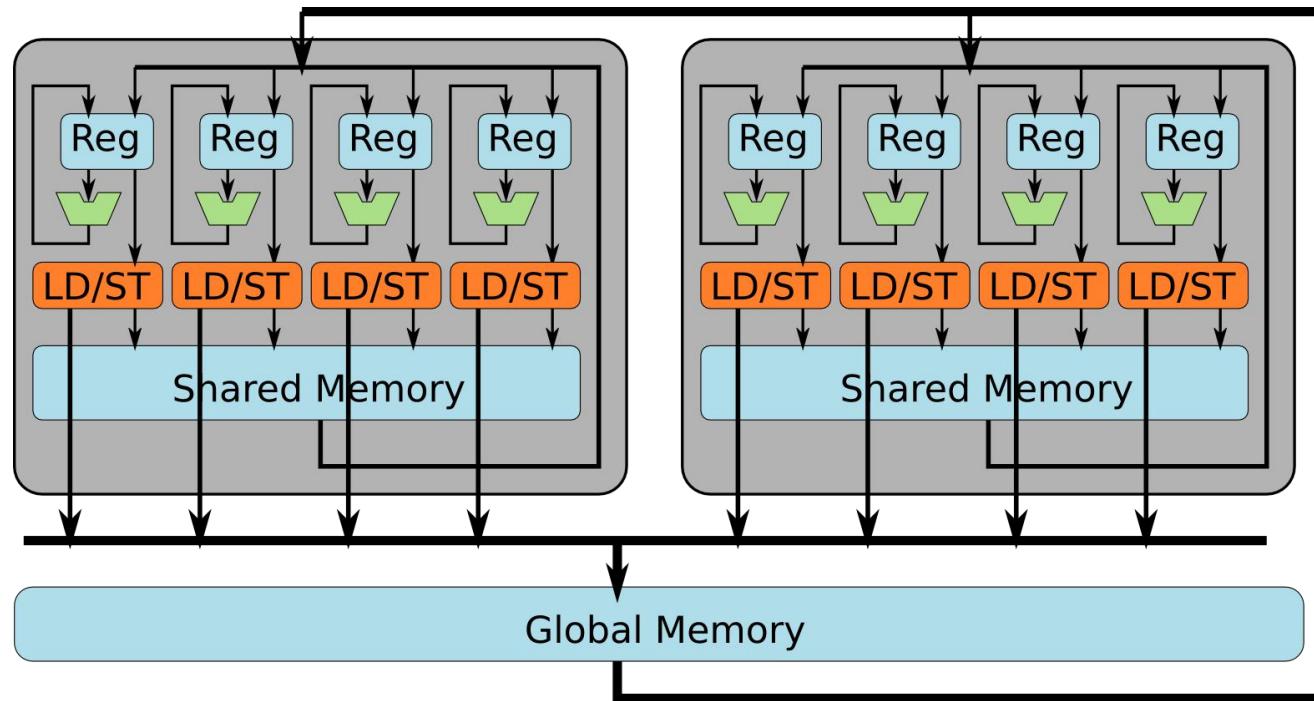
CUDA 程序员的 GPUs 视角

一个GPU具有多个SIMD单元.



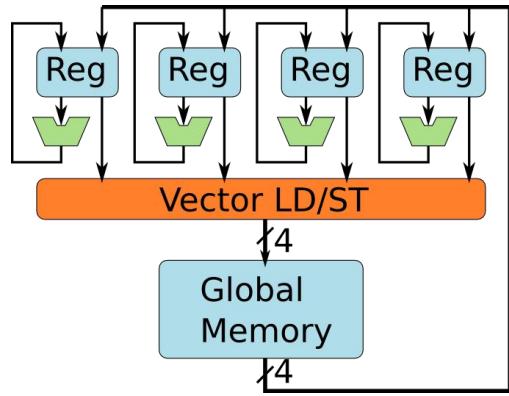
CUDA 程序员的 GPUs 视角 (2)

A GPU contains multiple SIMD Units. 所有 SIMD 单元都可以访问全局内存.

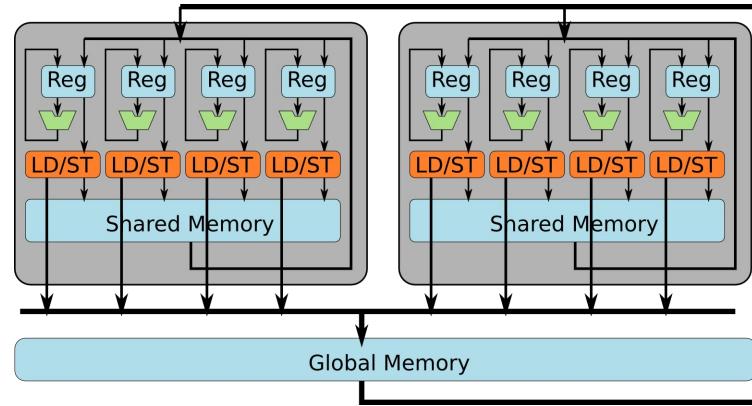


SIMD CPU(SSE)与GPU的区别

SSE



GPU

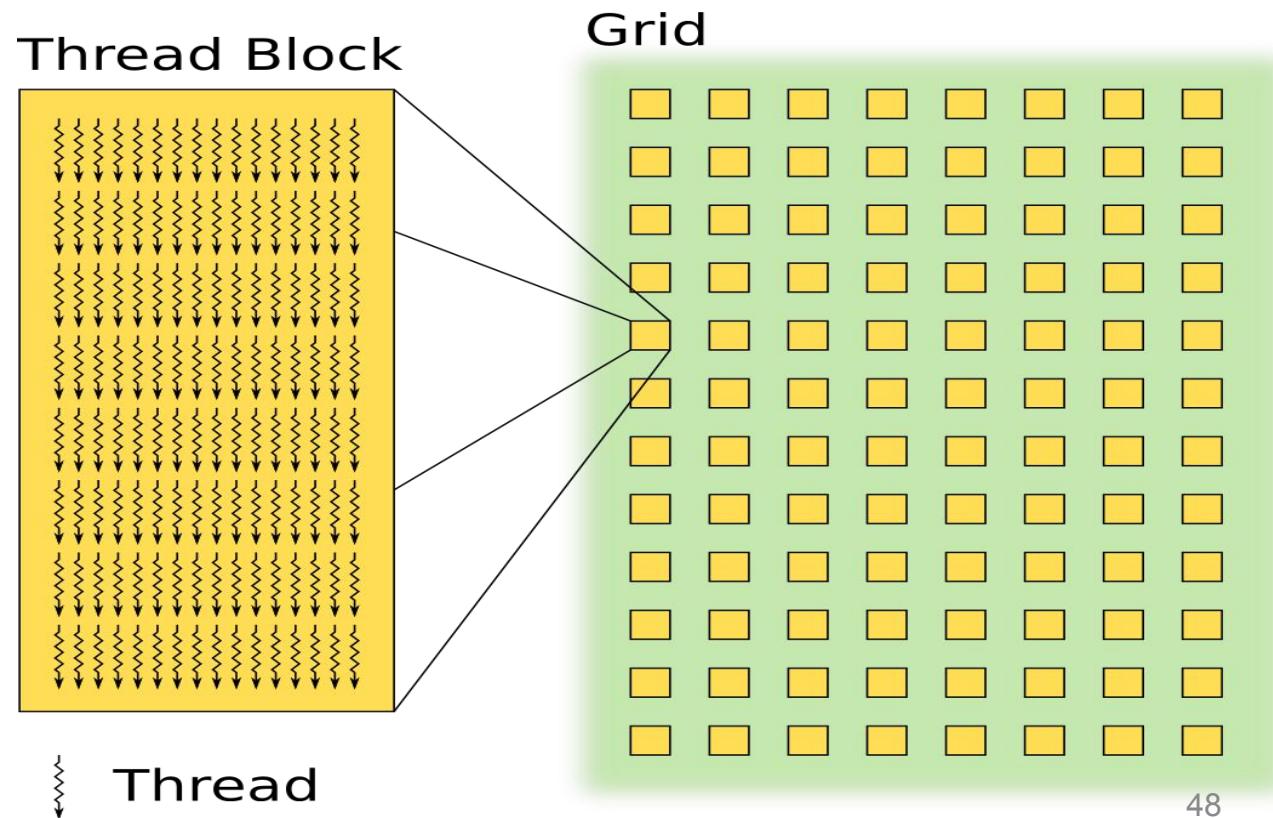


两个主要不同：

1. GPUs 使用线程（**threads**）代替矢量（**vectors**）
2. GPUs 具有共享内存（**Shared Memory**）空间

CUDA的线程层次结构

- 线程网格(Grid)
包含线程块
(Thread Blocks)
- 线程块包含
线程(Threads)



从C到CUDA C

转换成 CUDA

```
int A[2][4];
for(i=0;i<2;i++)
    for(j=0;j<4;j++)
        A[i][j]++;

```

被线程网格中所有线程执行的函数，在CUDA术语中，该函数称为Kernel. 即核函数

```
int A[2][4];
kernelF<<<(2,1),(4,1)>>>(A); // define 2x4=8 threads
__device__ kernelF(A){ // all threads run same kernel
    i = blockIdx.x; // each thread block has its id
    j = threadIdx.x; // each thread has its id
    A[i][j]++;
}
```

线程层次结构

例如：

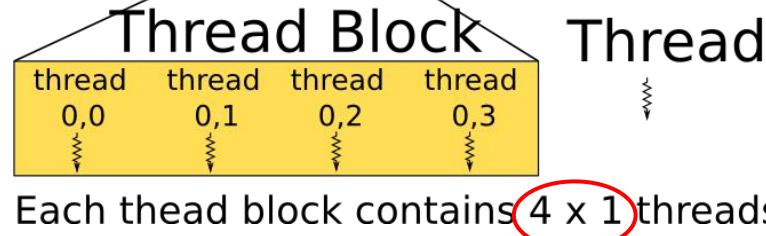
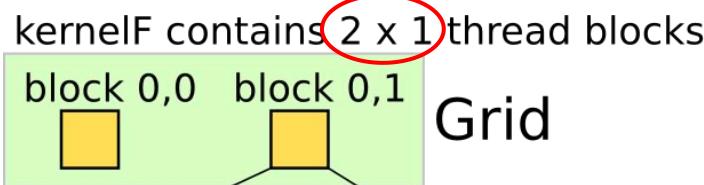
thread 3 of block 1 operates
on element A[1][3]

```
int A[2][4];
kernelF<<<(2,1),(4,1)>>>(A); // define 2x4=8 threads
__device__ kernelF(A){
    i = blockIdx.x;
    j = threadIdx.x;
    A[i][j]++;
}
```

Diagram illustrating the thread hierarchy:

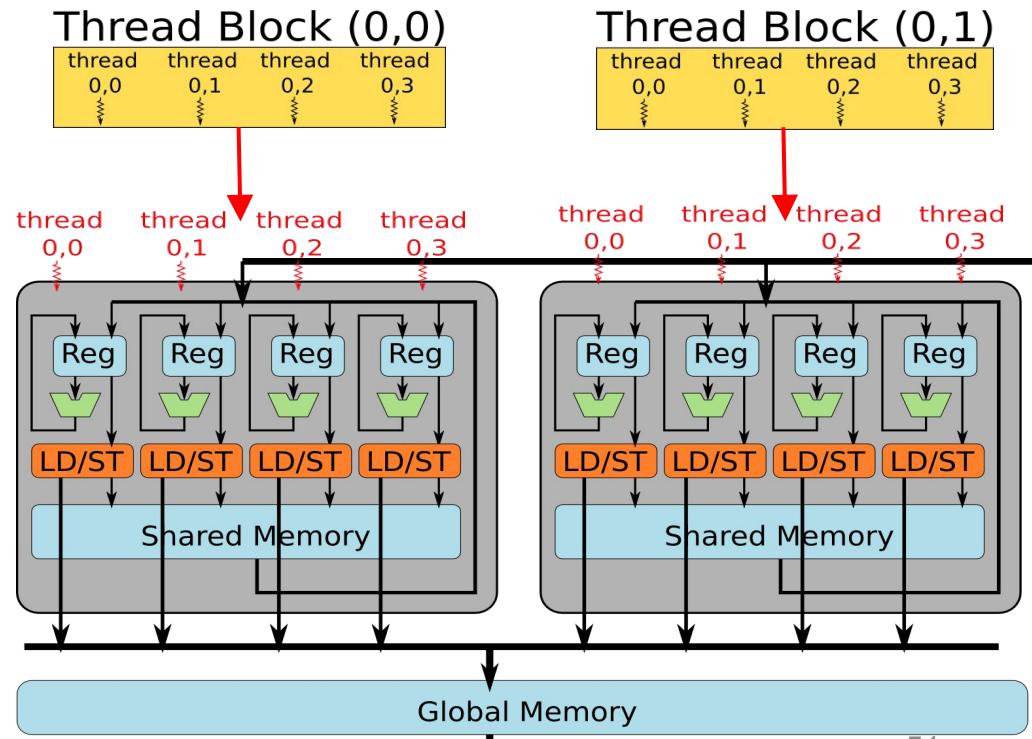
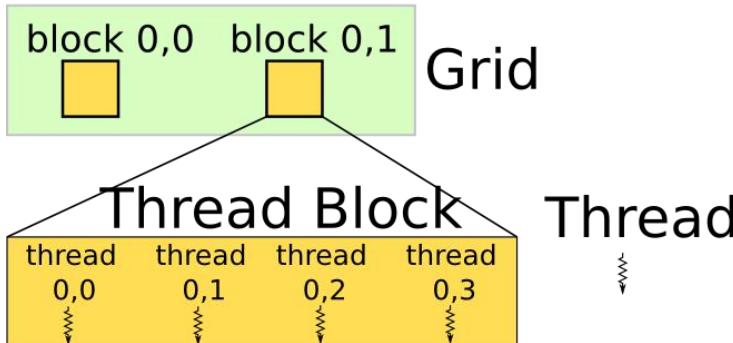
- Grid:** Contains 2 thread blocks (block 0,0 and block 0,1).
- Thread Block:** Contains 4 threads (thread 0,0, 0,1, 0,2, 0,3).
- Thread:** Each thread block contains 4 threads.

The diagram shows a hierarchical structure from Grid to Thread. A red arrow points from the text "thread 3 of block 1 operates on element A[1][3]" to the "Thread Block" section of the diagram.



线程分配

kernelF contains 2×1 thread blocks



线程块动态分配到SIMD组中

Grid

kernelF contains 2×2 thread blocks

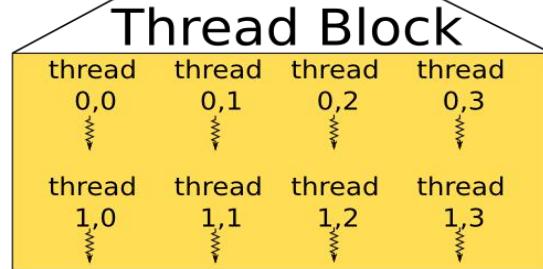
block 0,0 block 0,1



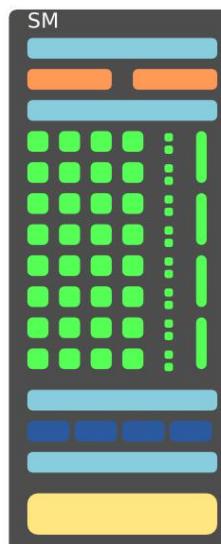
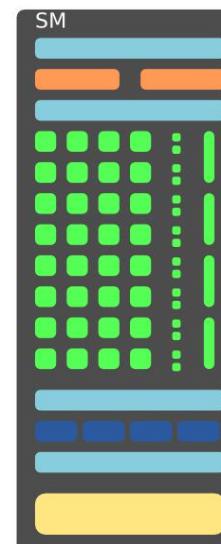
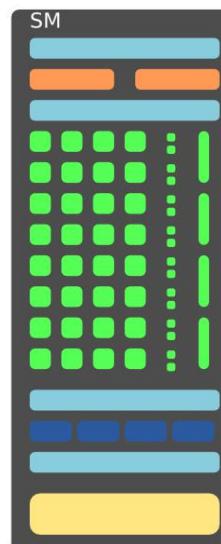
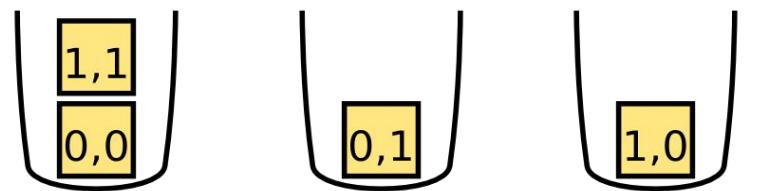
block 1,0 block 1,1



Thread



Each thread block contains 4×2 threads



线程执行

```
int A[2][4];
kernelF<<<(2,1),(4,1)>>>(A);
device kernelF(A){
    i = blockIdx.x;
    j = threadIdx.x;
    A[i][j]++;
}
```

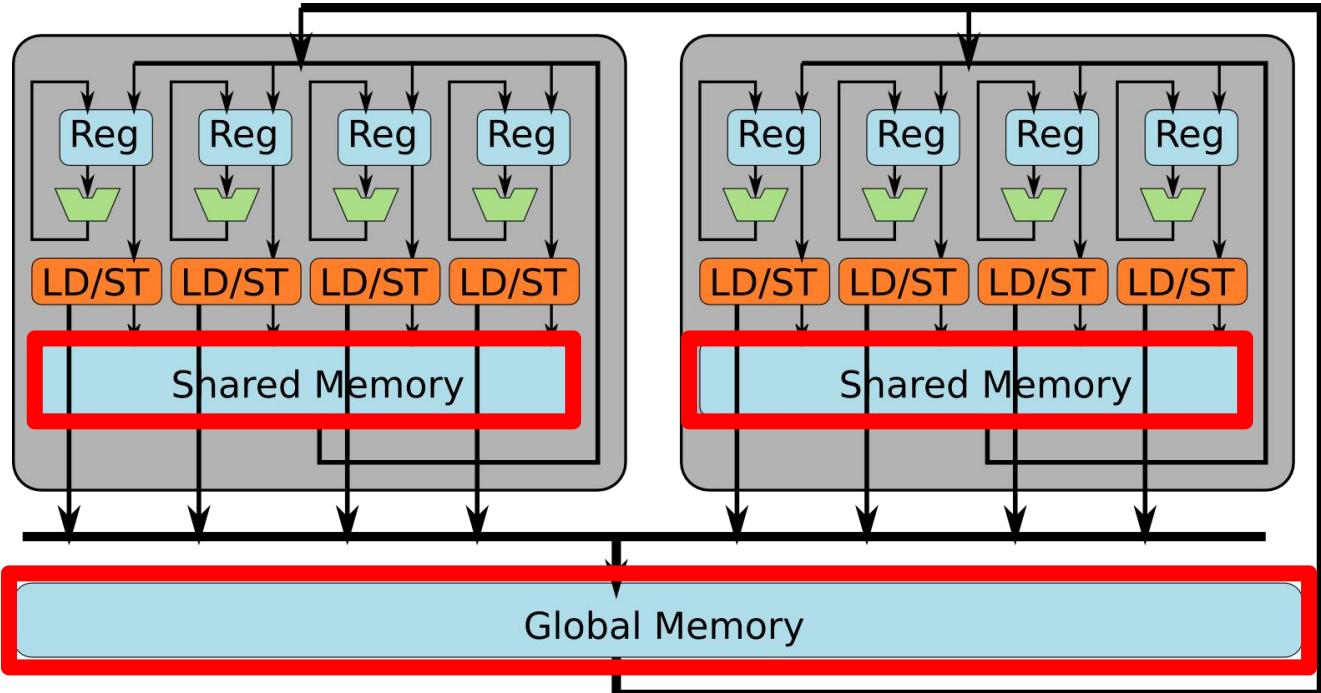
mv.u32	%r0, %ctaid.x	// r0 = i = blockIdx.x
mv.u32	%r1, %ntid.x	// r1 = "threads-per-block"
mv.u32	%r2, %tid.x	// r2 = j = threadIdx.x
mad.u32	%r3, %r2, %r1, %r0	// r3 = i * "threads-per-block" + j
ld.global.s32	%r4, [%r3]	// r4 = A[i][j]
add.s32	%r4, %r4, 1	// r4 = r4 + 1
st.global.s32	[%r3], %r4	// A[i][j] = r4

利用内存层次结构

内存访问
延迟

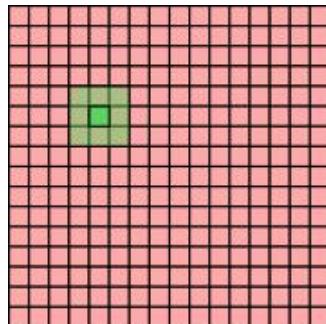
several cycles

100+ cycles



示例: 均值滤波

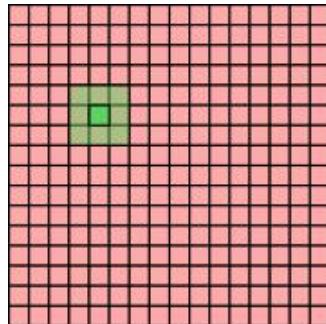
Average over a
3x3 window for
a 16x16 array



```
kernelF<<<(1,1),(16,16)>>>(A);  
__device__ kernelF(A){  
    i = threadIdx.y;  
    j = threadIdx.x;  
    tmp = (A[i-1][j-1] + A[i-1][j] +  
           ... + A[i+1][j+1] ) / 9;  
    A[i][j] = tmp;  
}  
Each thread loads 9 elements  
from global memory.  
It takes hundreds of cycles.
```

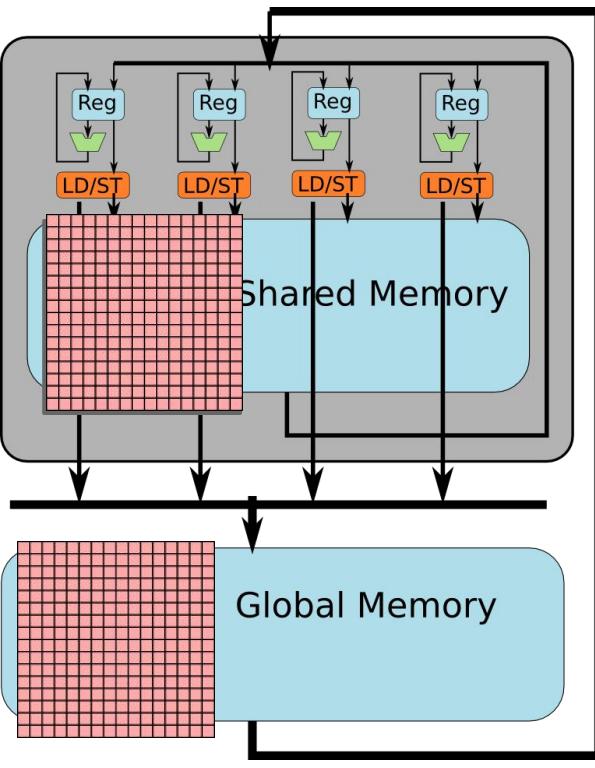
利用共享内存

Average over a
3x3 window for
a 16x16 array



```
kernelF<<<(1,1),(16,16)>>>(A);  
__device__ kernelF(A){  
    __shared__ int smem[16][16];  
    i = threadIdx.y;  
    j = threadIdx.x;  
    smem[i][j] = A[i][j]; // load to smem  
    A[i][j] = ( smem[i-1][j-1] + smem[i-1][j] +  
                ... + smem[i+1][i+1] ) / 9;  
}
```

利用共享内存 (2)



```
kernelF<<<(1,1),(16,16)>>>(A);  
__device__ kernelF(A){  
    __shared__ smem[16][16];  
    i = threadIdx.y;  
    j = threadIdx.x;    Each thread loads one  
    smem[i][j] = A[i][j]; // load to smem  
    A[i][j] = ( smem[i-1][j-1] + smem[i-1][j] +  
                ... + smem[i+1][j+1] ) / 9;  
}
```

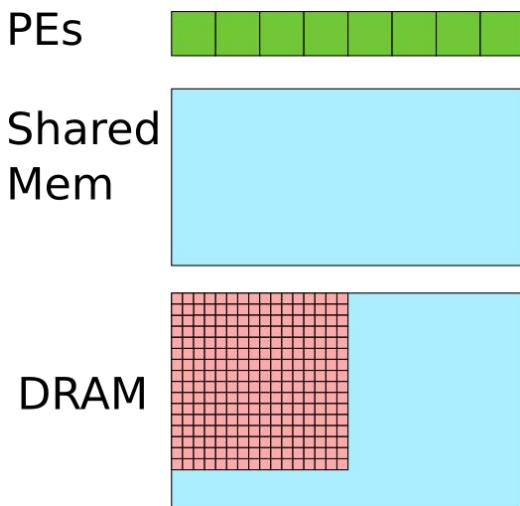
但是， 程序出错了！

Hazards!

```
kernelF<<<(1,1),(16,16)>>>(A);  
__device__  kernelF(A){  
    __shared__ smem[16][16];  
    i = threadIdx.y;  
    j = threadIdx.x;  
    smem[i][j] = A[i][j]; // load to smem  
    A[i][j] = ( smem[i-1][j-1] + smem[i-1][j] +  
                ... + smem[i+1][j+1] ) / 9;  
}
```

为什么会出现错误？

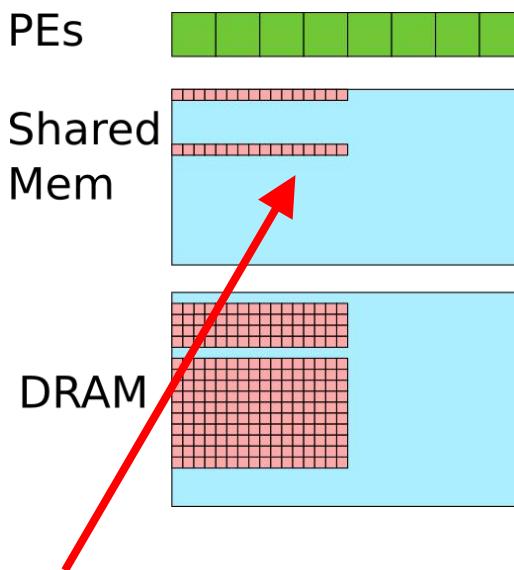
Assume 256 threads are scheduled on 8 PEs.



```
kernelF<<<(1,1),(16,16)>>>(A);  
__device__ kernelF(A){  
    __shared__ smem[16][16];  
    i = threadIdx.y;  
    j = threadIdx.x;      Before load instruction  
    smem[i][j] = A[i][j]; // load to smem  
    A[i][j] = ( smem[i-1][j-1] + smem[i-1][j] +  
                ... + smem[i+1][j+1] ) / 9;  
}
```

为什么会出现错误？（2）

Assume 256 threads are scheduled on 8 PEs.

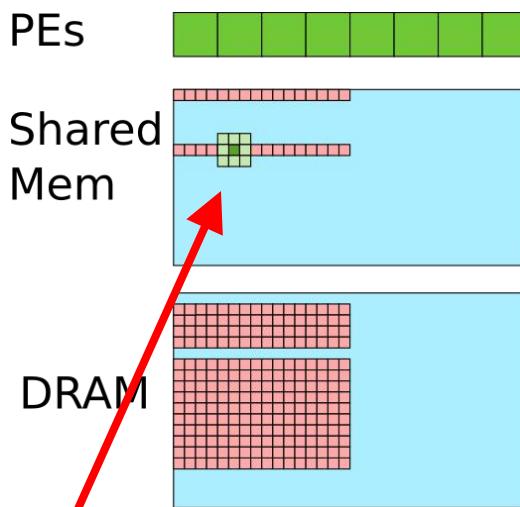


Some threads finish the load earlier than others.

```
kernelF<<<(1,1),(16,16)>>>(A);  
__device__ kernelF(A){  
    __shared__ smem[16][16];  
    i = threadIdx.y;  
    j = threadIdx.x;  
    smem[i][j] = A[i][j]; // load to smem  
    A[i][j] = ( smem[i-1][j-1] + smem[i-1][j] +  
                ... + smem[i+1][j+1] ) / 9;  
}
```

为什么会出现错误？（3）

Assume 256 threads are scheduled on 8 PEs.

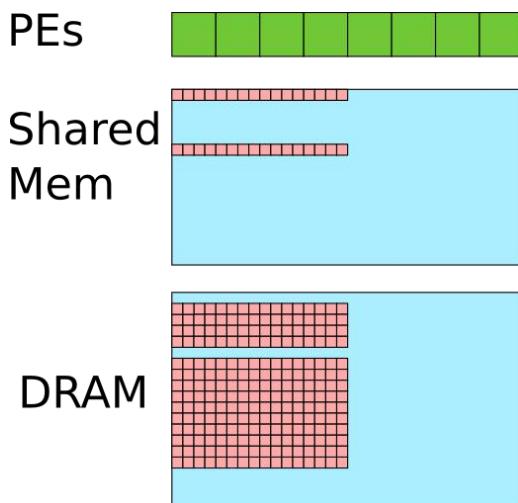


Some elements in the window
are not yet loaded by other
threads. Error!

```
kernelF<<<(1,1),(16,16)>>>(A);  
__device__ kernelF(A){  
    __shared__ smem[16][16];  
    i = threadIdx.y;  
    j = threadIdx.x;  
    smem[i][j] = A[i][j]; // load to smem  
    A[i][j] = ( smem[i-1][j-1] + smem[i-1][j] +  
                ... + smem[i+1][j+1] ) / 9;  
}
```

解决办法？

Assume 256 threads are
scheduled on 8 PEs.

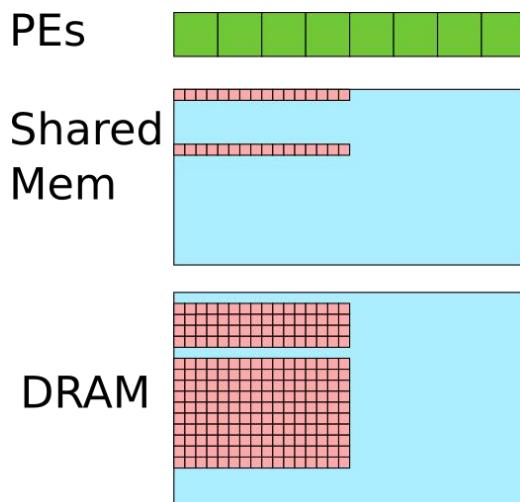


```
kernelF<<<(1,1),(16,16)>>>(A);  
__device__ kernelF(A){  
    __shared__ smem[16][16];  
    i = threadIdx.y;  
    j = threadIdx.x;  
    smem[i][j] = A[i][j]; // load to smem  
    A[i][j] = ( smem[i-1][j-1] + smem[i-1][j] +  
                ... + smem[i+1][j+1] ) / 9;  
}
```

使用 "SYNC" 同步障

kernelF<<<(1,1),(16,16)>>>(A);

Assume 256 threads are
scheduled on 8 PEs.



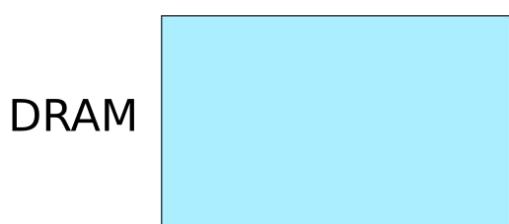
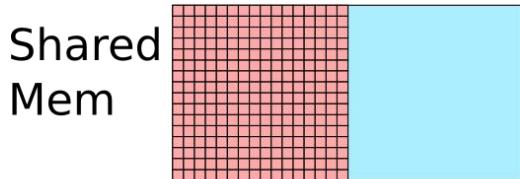
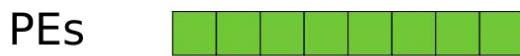
```
__device__ kernelF(A){  
    __shared__ smem[16][16];  
    i = threadIdx.y;  
    j = threadIdx.x;  
    smem[i][j] = A[i][j]; // load to smem  
    __SYNC();  
    A[i][j] = ( smem[i-1][j-1] + smem[i-1][j] +  
                ... + smem[i+1][j+1] ) / 9;
```

}

使用 "SYNC" 同步障 (2)

kernelF<<<(1,1),(16,16)>>>(A);

Assume 256 threads are
scheduled on 8 PEs.



device kernelF(A){

shared smem[16][16];

i = threadIdx.y;

j = threadIdx.x;

smem[i][j] = A[i][j]; // load to smem

SYNC();

A[i][j] = (smem[i-1][j-1] + smem[i-1][j] +

... + smem[i+1][i+1]) / 9;

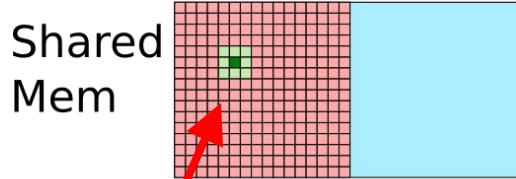
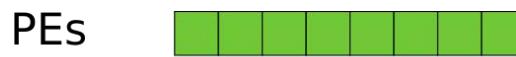
Wait until all
threads
hit barrier

}

使用 "SYNC" 同步障 (3)

kernelF<<<(1,1),(16,16)>>>(A);

Assume 256 threads are
scheduled on 8 PEs.



DRAM

device kernelF(A){

shared smem[16][16];

i = threadIdx.y;

j = threadIdx.x;

smem[i][j] = A[i][j]; // load to smem

SYNC();

A[i][j] = (smem[i-1][j-1] + smem[i-1][j] +
... + smem[i+1][j+1]) / 9;

}

All elements in the window
are loaded when each
thread starts averaging.

目录

- GPU微处理器简介
- CUDA编程
 - 引言：开发平台、异构计算
 - CUDA编程模型：线程、线程块、线程网格
 - CUDA执行模型：流多处理器和线程束
 - 内存层次结构
 - 编程实例

CUDA执行模型：流多处理器

- **SP**: GPU内核上运行**kernel**函数时，相同的指令序列会被大量称为流处理器（**SP**）的处理单元同步执行。
 - 每个线程运行在一个SP上， SP称为内核（CUDA Core）
- **SM**: 同一控制单元控制下执行的一组**SP**称为流多处理器（**SM**）。
 - GPU可以包含多个SM， 每个运行自己的**kernel**函数。

CUDA执行模型：SM和warp

- 执行模型会提供一个操作视图，说明如何在特定的计算架构上执行指令。
- CUDA采用单指令多线程（SIMT）架构来管理和执行线程，每32个（大小与硬件相关）线程为一组，被称为线程束（warp）。
 - 线程束中的所有线程同时执行相同的指令。
 - 每个SM都将分配给它的线程块分到线程束中，然后在可用的硬件资源上调度执行。

讨论:

1. 单指令多线程 (SIMT)
2. 共享内存

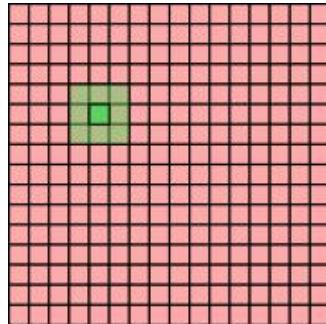
Q: 显式管理内存的优
缺点是什么?



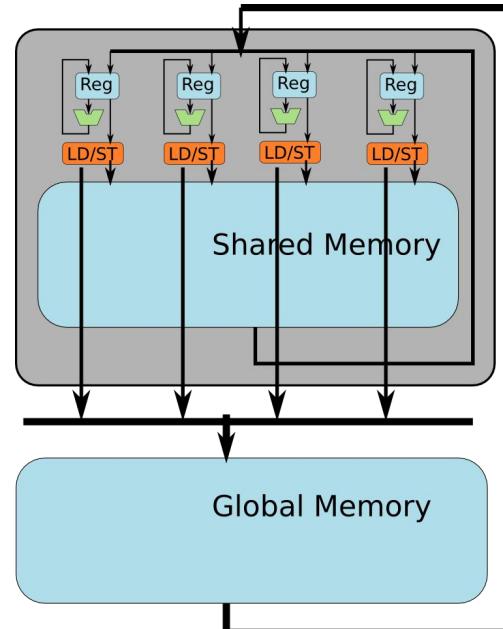
Q: SIMT 与 vector SIMD 编程模型的区别是什么?

回顾前一个示例

Average over a
3x3 window for
a 16x16 array



假设 vector SIMD 和 SIMT 都有共
享内存.
它们的区别?



Vector SIMD v.s. SIMT

```
int A[16][16]; // A in global memory
__shared__ int B[16][16]; // B in shared mem

for(i=0;i<16;i++){
    for(j=0;j<16;j+=4){
        movups xmm0, [ &A[i][j] ]
        movups [ &B[i][j] ], xmm0 }
    for(i=0;i<16;i++){
        for(j=0;j<16;j+=4){
            addps xmm1, [ &B[i-1][j-1] ]
            addps xmm1, [ &B[i-1][j] ]
            ... divps xmm1, 9 }
        for(i=0;i<16;i++){
            for(j=0;j<16;j+=4){
                addps [ &A[i][j] ], xmm1 }}}
```



```
kernelF<<<(1,1),(16,16)>>>(A);

__device__ kernelF(A){

    __shared__ int smem[16][16];
    i = threadIdx.y;
    j = threadIdx.x;
    smem[i][j] = A[i][j]; // load to smem (1)
    __sync(); // threads wait at barrier
    A[i][j] = ( smem[i-1][j-1] + smem[i-1][j] +
    (3)      ... + smem[i+1][j+1] ) / 9;
}
```



- (1) load to shared mem
- (2) compute
- (3) store to global mem

Vector SIMD v.s. SIMT

```
int A[16][16]; // A in global memory
__shared__ int B[16][16]; // B in shared mem
for(i=0;i<16;i++){
    for(j=0;j<16;j+=4){ ← (a)
        movups xmm0, [ &A[i][j] ]
        movups [ &B[i][j] ], xmm0 } } ↑ (c)
for(i=0;i<16;i++){
    for(j=0;j<16;j+=4){
        addps xmm1, [ &B[i-1][j-1] ]
        addps xmm1, [ &B[i-1][j] ]
        ... divps xmm1, 9 } }
for(i=0;i<16;i++){
    for(j=0;j<16;j+=4){
        addps [ &A[i][j] ], xmm1 } }
```

```
kernelF<<<(1,1),(16,16)>>>(A);
__device__ kernelF(A){
    __shared__ smem[16][16];
    i = threadIdx.y; ↑ (b)
    j = threadIdx.x; ↓ (b)
    smem[i][j] = A[i][j]; // load to smem
    __sync(); // threads wait at barrier ← (d)
    A[i][j] = ( smem[i-1][j-1] + smem[i-1][j] +
                ... + smem[i+1][j+1] ) / 9;
}
```

- (a) HW vector width explicit to programmer
- (b) HW vector width transparent to programmers
- (c) each vector executed by all PEs in lock step
- (d) threads executed out of order, need explicit sync

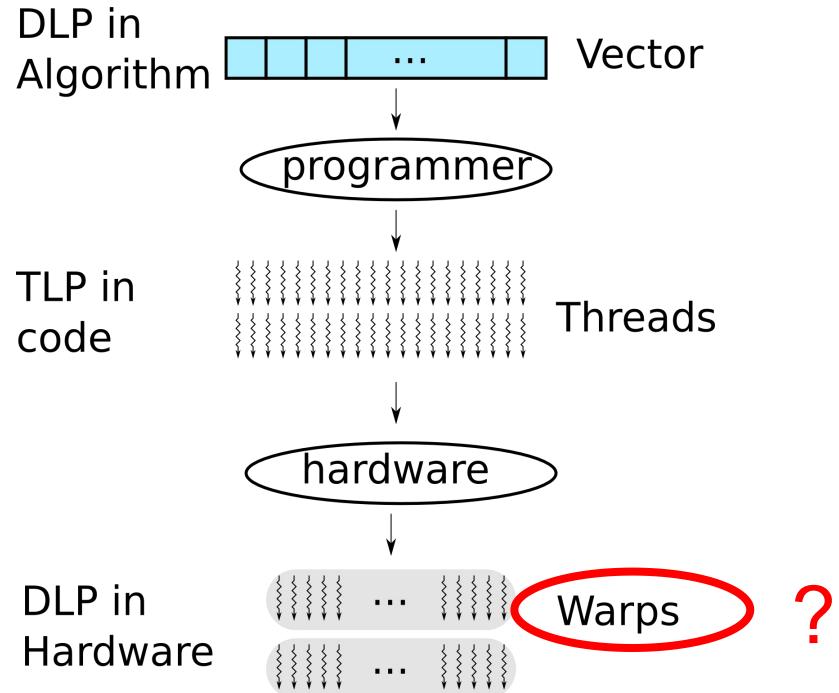
SIMT vs. SIMD

- 相同点：两者都是将相同的指令广播给多个执行单元来实现并行。
- 不同点：
 - SIMD要求同一个向量中的所有元素要在一个统一的同步组中一起执行；
 - SIMT允许属于同一线程束的多个线程独立执行。
- SIMT包含3个SIMD不具备的关键特征：
 - 每个线程都有自己的指令地址计数器
 - 每个线程都有自己的寄存器状态
 - 每个线程可以有一个独立的执行路径

SIMT : SM和线程块

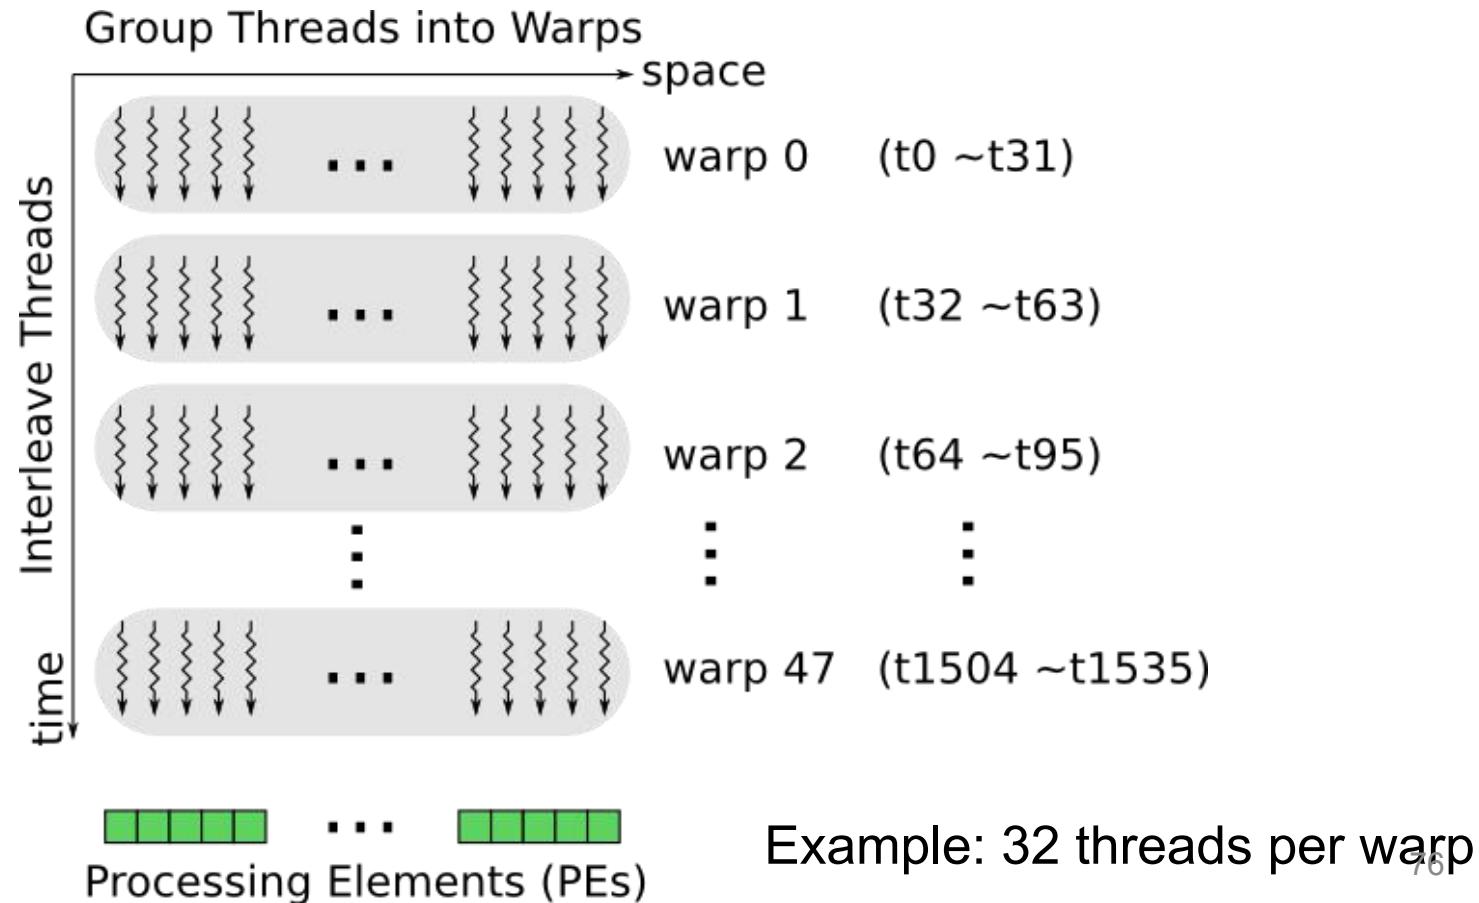
- 一个**线程块**只能在一个**SM**上被调度
- 一个**SM**可以容纳多个线程块
- **寄存器**和**共享内存**是**SM**中的稀缺资源
- 线程块中的不同线程可能会以**不同的速度前进**。

回顾：数据级并行到线程级并行



Programmers convert **data level parallelism (DLP)** into **thread level parallelism (TLP)**.

硬件 (HW) 将线程分为不同的线程束 (Warp)

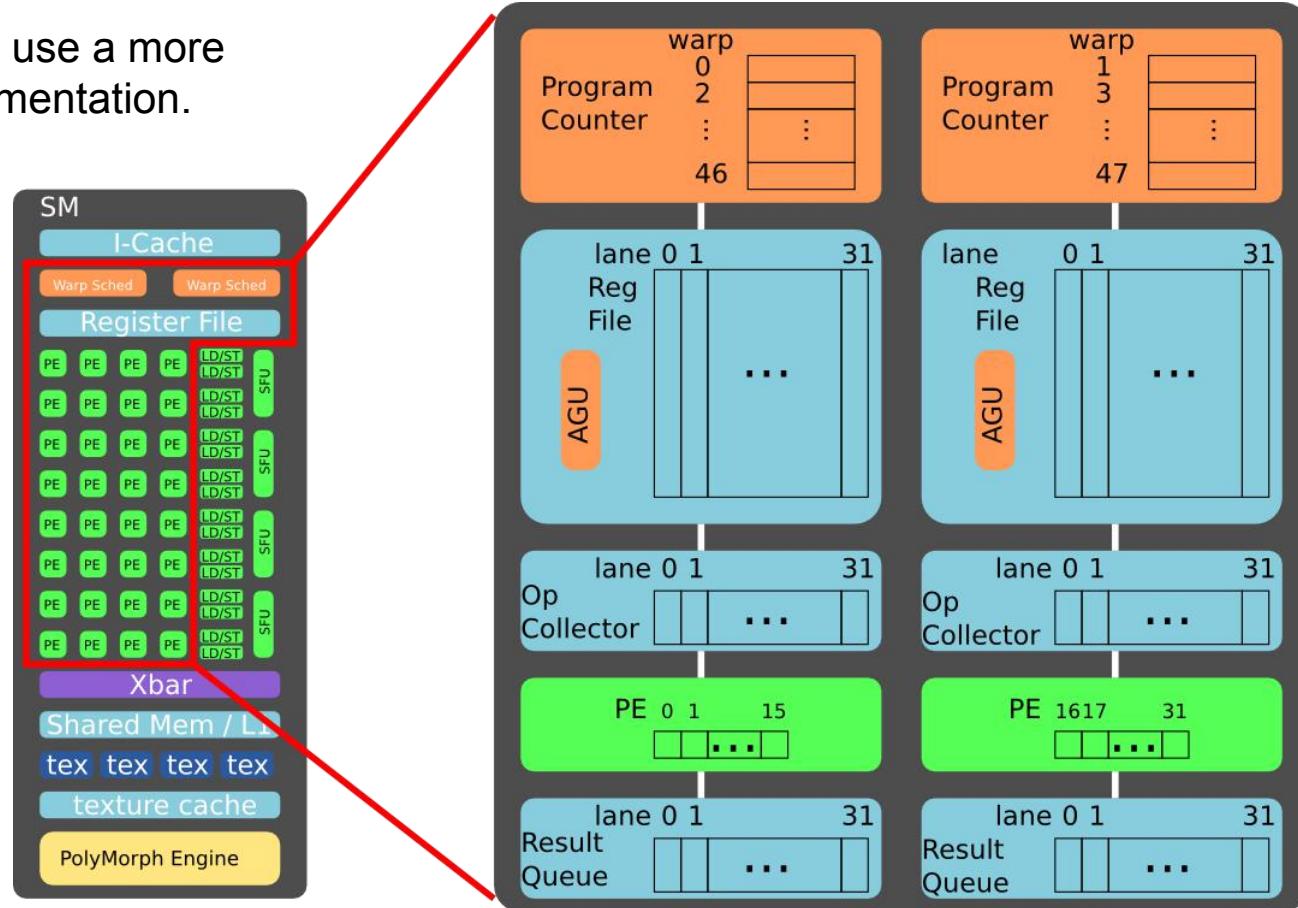


线程和线程束的执行

举例

举例：硬件结构

Note: NVIDIA may use a more complicated implementation.



举例：寄存器分配

Assumption:

register file has 32 lanes

each warp has 32 threads

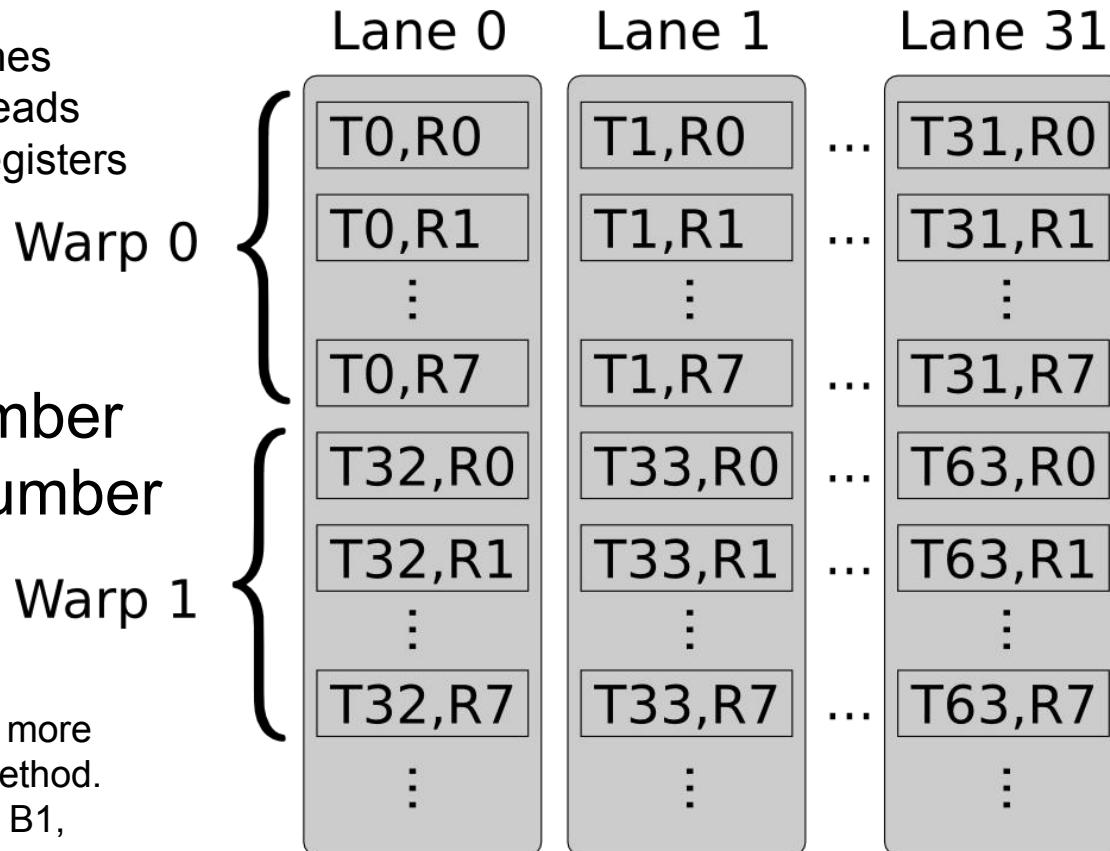
each thread uses 8 registers

Acronym:

“T”: thread number

“R”: register number

Note: NVIDIA may use a more complicated allocation method.
See patent: US 7634621 B1,
“Register file allocation”.



举例：线程束

Address : Instruction

0x0004 : add r0, r1, r2

0x0008 : sub r3, r4, r5

假设：

两个数据通道

warp 0 在左边的数据通道

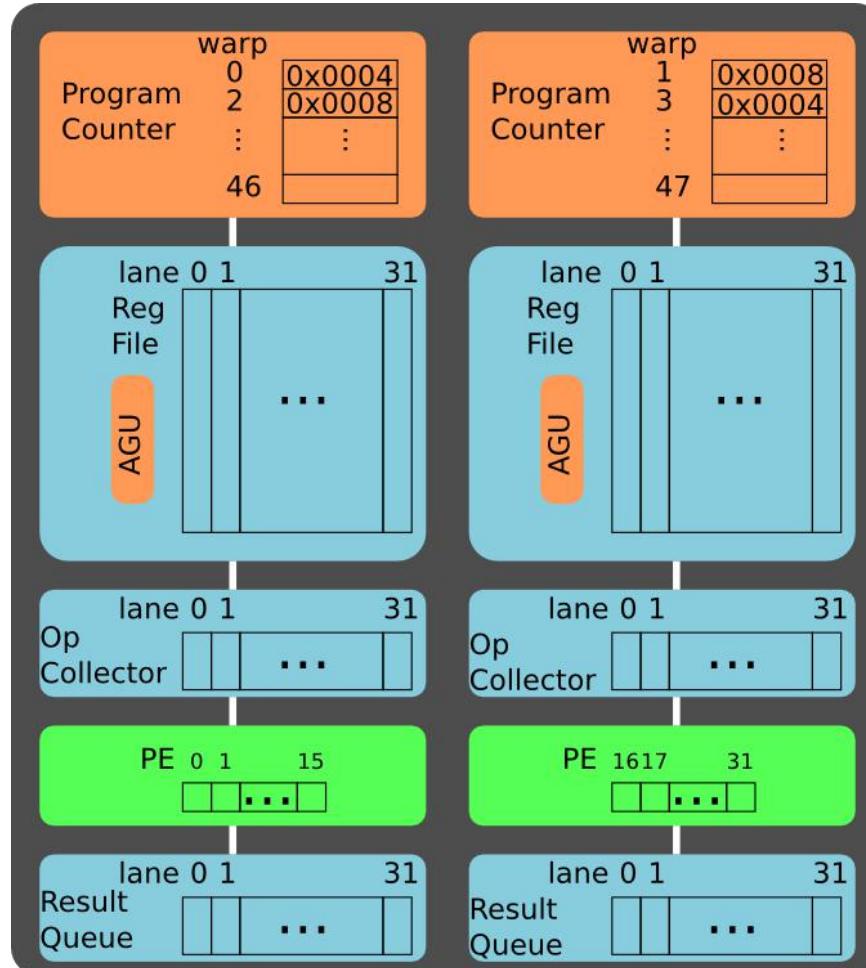
warp 1 在右边的数据通道

术语：

“**AGU**”: address generation unit

“r”: register in a thread

“w”: warp number



Read Src Op 1

Address : Instruction

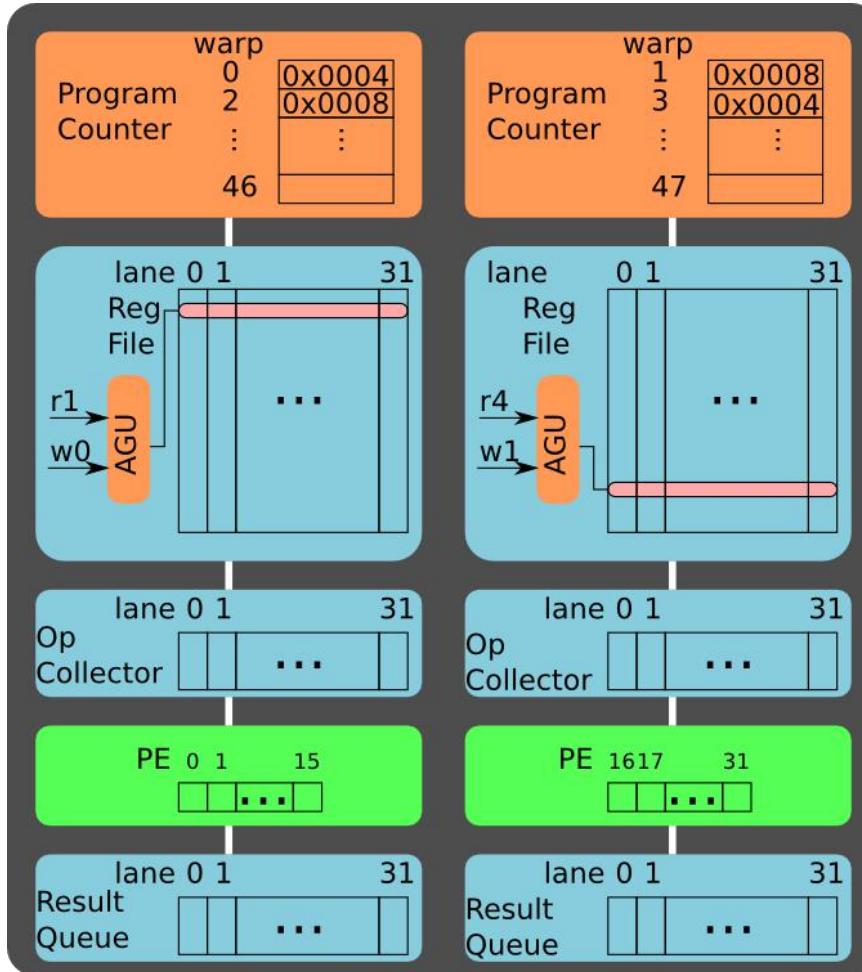
0x0004 : add r0, **r1**, r2

0x0008 : sub r3, **r4**, r5

Read source operands:

r1 for warp 0

r4 for warp 1



Buffer Src Op 1

Address : Instruction

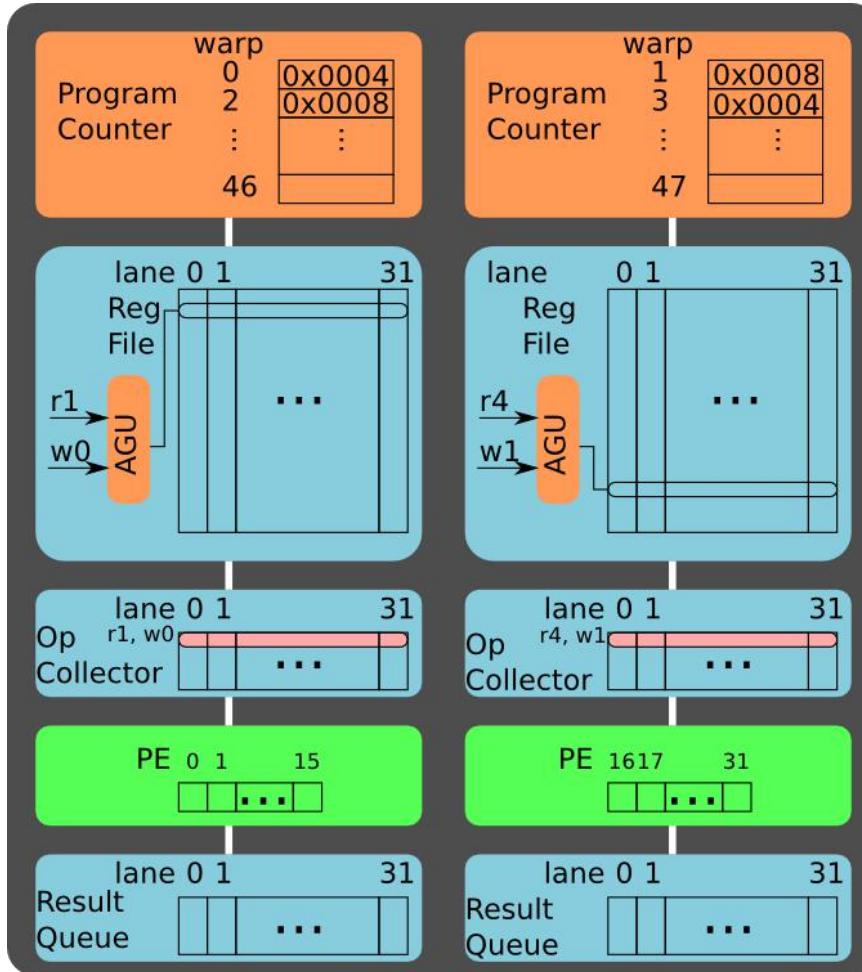
0x0004 : add r0, **r1**, r2

0x0008 : sub r3, **r4**, r5

Push ops to op collector:

r1 for warp 0

r4 for warp 1



Read Src Op 2

Address : Instruction

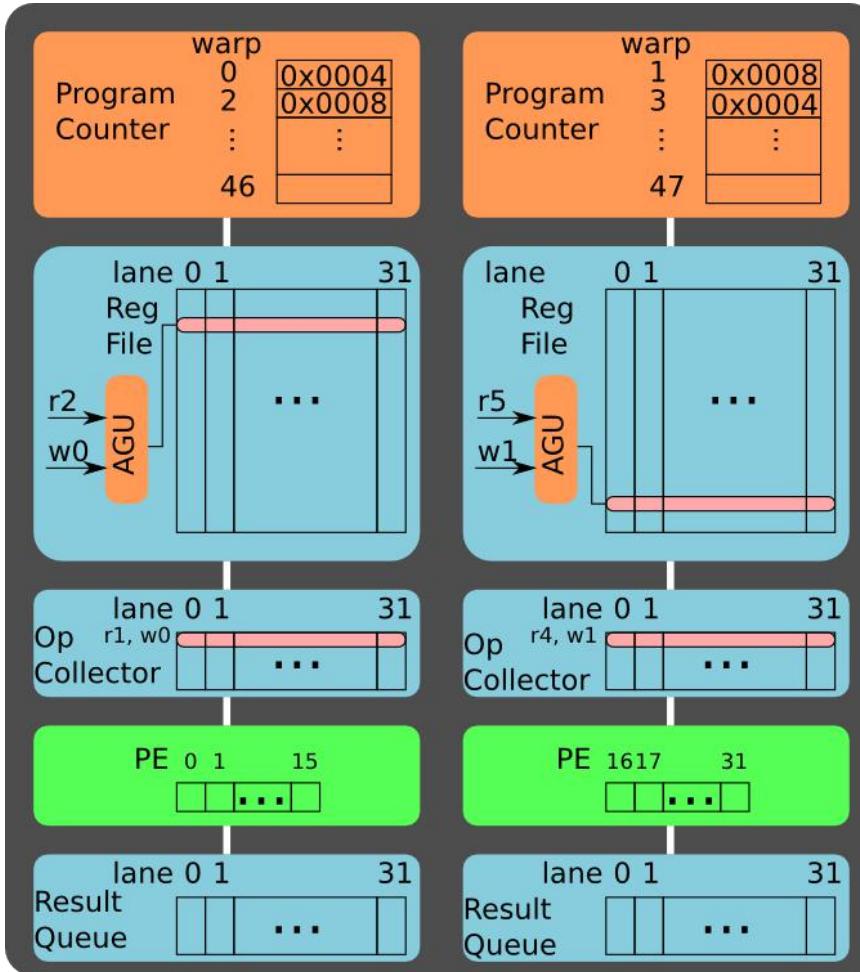
0x0004 : add r0, r1, **r2**

0x0008 : sub r3, r4, **r5**

Read source operands:

r2 for warp 0

r5 for warp 1

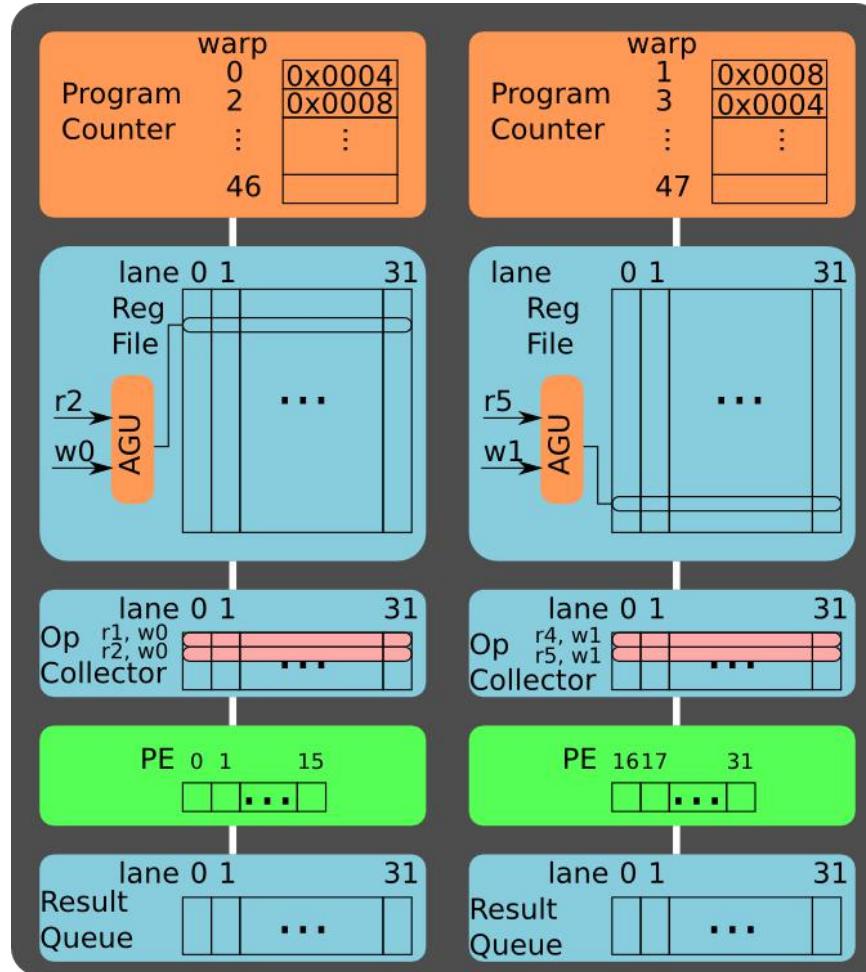


Buffer Src Op 2

Address : Instruction

0x0004 : add r0, r1, **r2**

0x0008 : sub r3, r4, **r5**



Push ops to op collector:

r2 for warp 0

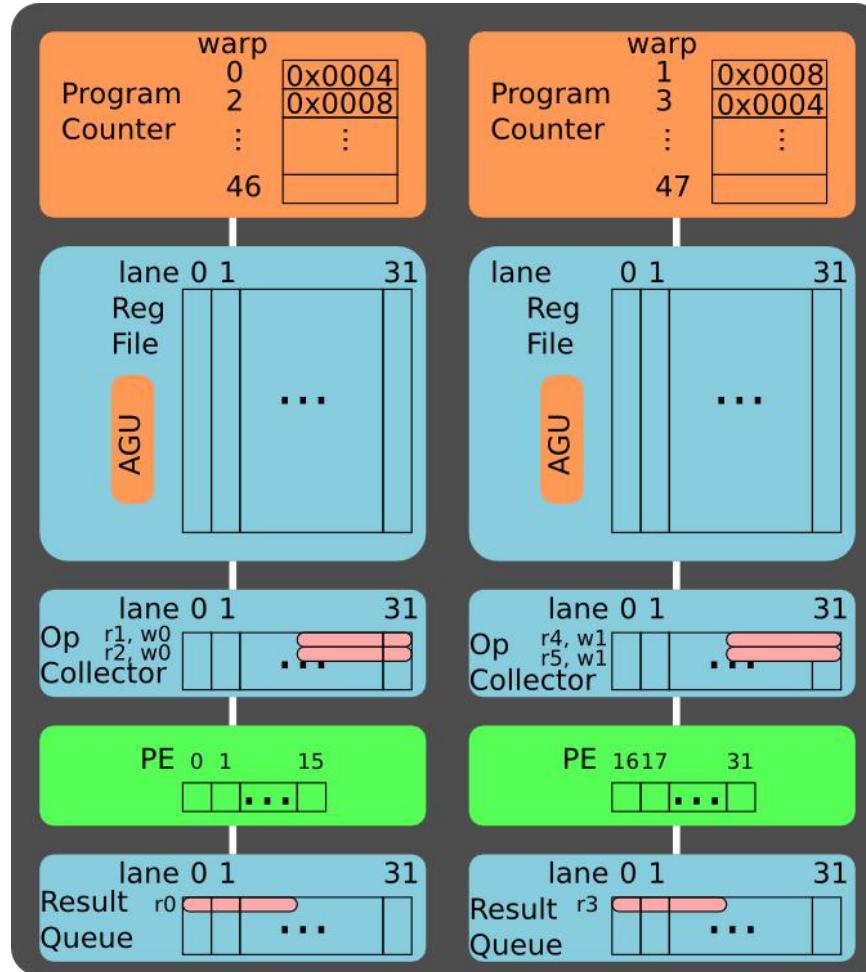
r5 for warp 1

Execute Stage 1

Address : Instruction

0x0004 : **add** r0, r1, r2

0x0008 : **sub** r3, r4, r5



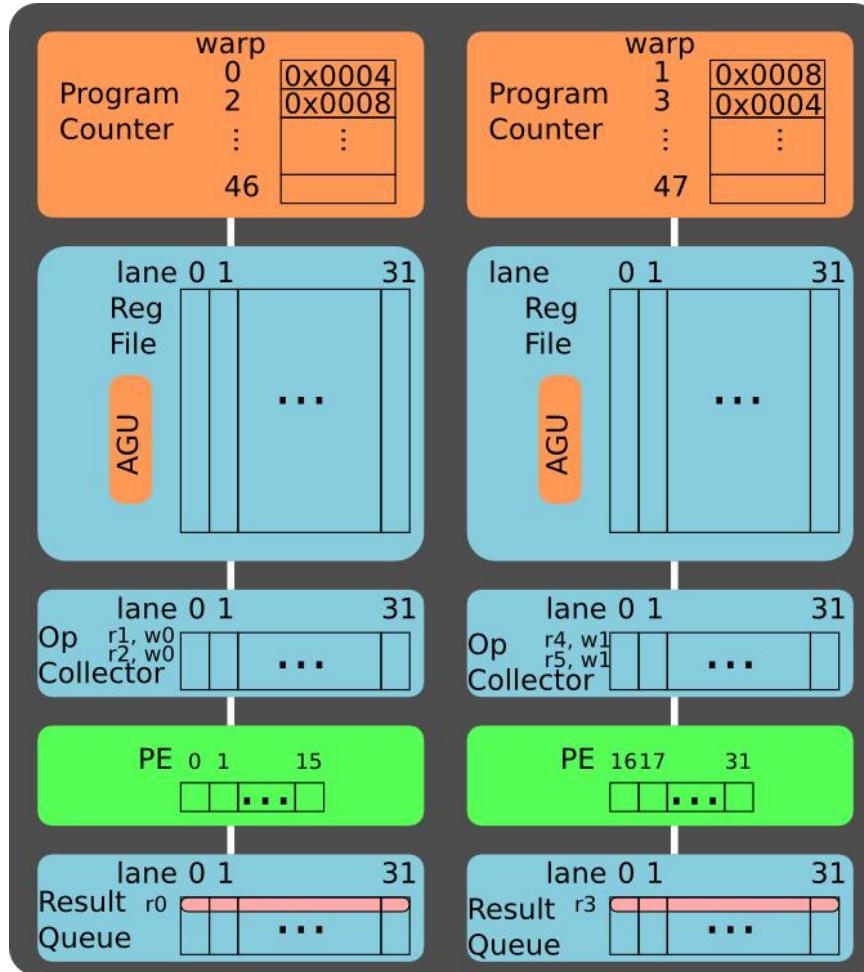
Compute the **first 16 threads** in the warp.

Execute Stage 2

Address : Instruction

0x0004 : **add** r0, r1, r2

0x0008 : **sub** r3, r4, r5



Compute the **last 16 threads** in the warp.

Write Back

Address : Instruction

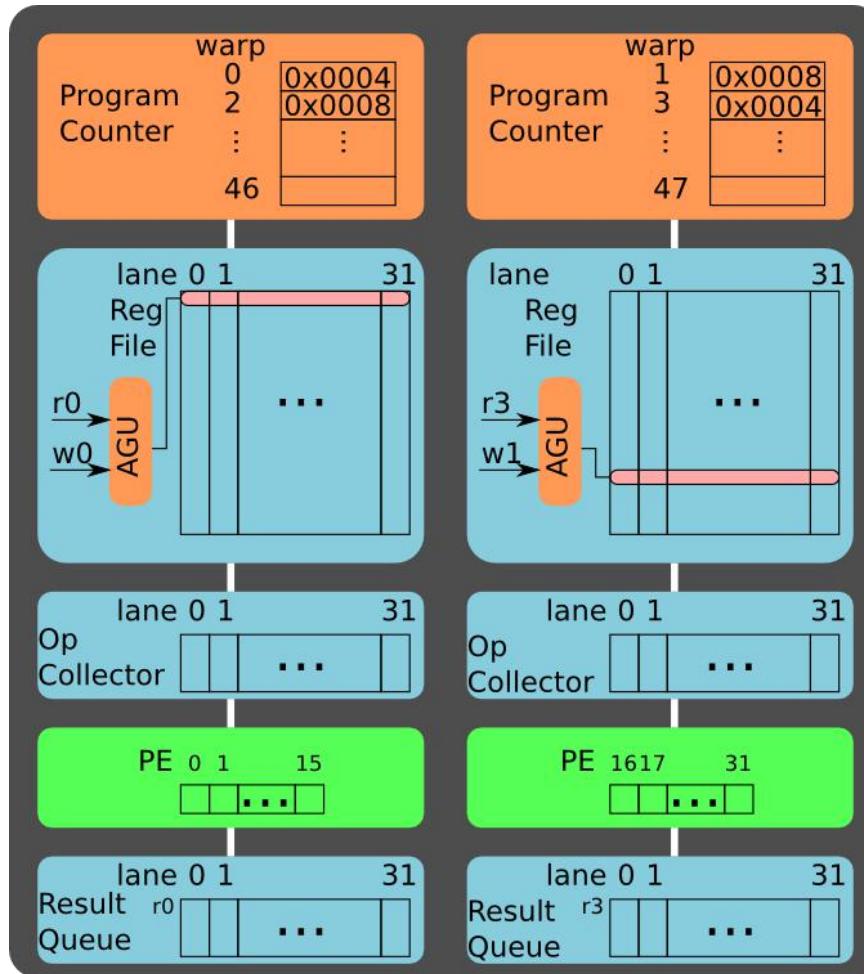
0x0004 : add r0, r1, r2

0x0008 : sub r3, r4, r5

Write back:

r0 for warp 0

r3 for warp 1



目录

- GPU微处理器简介
- CUDA编程
 - 引言：开发平台、异构计算
 - CUDA编程模型：线程、线程块、线程网格
 - CUDA执行模型：流多处理器和线程束
 - 内存层次结构
 - 编程实例

CUDA存储模型

- 对程序员来说，一般有两类存储器：
 - 可编程的：显式控制哪些数据存放在可编程内存中；
 - 不可编程的：程序员不能决定数据的存放位置，程序将自动生成存放位置以获得良好性能。（如CPU缓存）
- CUDA内存模型有多种可编程内存
 - 寄存器
 - 共享内存
 - 本地内存
 - 常量内存
 - 纹理内存
 - 全局内存

CUDA存储模型-寄存器

- 寄存器是GPU上运行速度最快的存储空间。
- 寄存器变量对于每个线程都是私有的。
- 寄存器变量与核函数的生命周期相同。
- Fermi GPU每个线程最多有63个寄存器，Kepler GPU每个线程最多255个寄存器
- 一个核函数使用了超过硬件限制数量的寄存器，则会用本地内存代替。

CUDA存储模型-本地内存、共享内存

- 本地内存：
 - 核函数中的线程都有自己私有的本地内存。
- 共享内存：
 - 线程块有自己的共享内存，对同一线程块中的所有线程都可见，是线程间相互通信的基本方式。
 - 共享内存是片上内存，类似CPU一级缓存，但可编程。
 - SM中的一级缓存和共享内存共同使用64KB的片上内存，可静态划分或运行时动态配置。

CUDA存储模型-常量内存、纹理内存、全局内存

➤ 常量内存：所有线程只读访问

- 驻留在设备内存中，并在每个SM专用的常量缓存中缓存。
- 常量内存是静态声明的，并对同一编译单元中的所有核函数可见。

➤ 纹理内存：所有线程只读访问

- 驻留在设备内存中，并在每个SM的只读缓存中缓存。
- 是一种通过指定的只读缓存访问的全局内存。

➤ 全局内存：所有线程都可以访问

- 是GPU中最大、延迟最高并且最常使用的内存。
- 全局内存变量可以被静态声明或动态声明。
- 全局内存常驻于设备内存中，可通过内存事务进行访问。

CUDA存储模型-GPU缓存

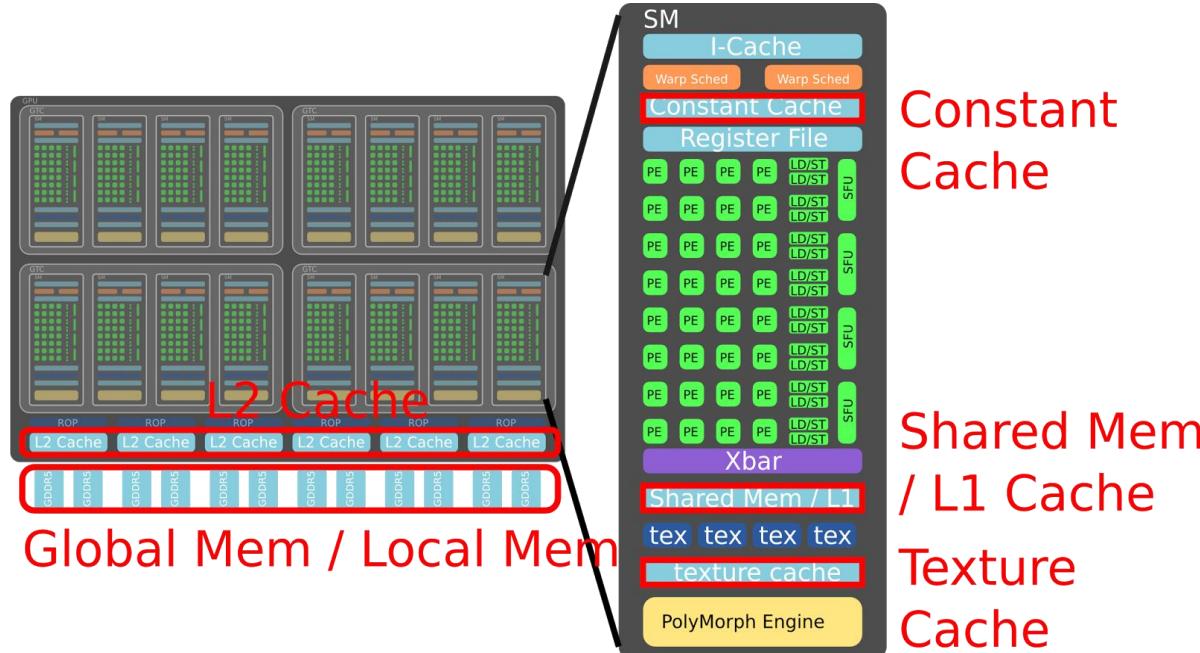
➤ GPU缓存是不可编程的内存。

- 一级缓存（每个SM有一个）
- 二级缓存（所有SM共享一个）
- 只读常量缓存
- 只读纹理缓存

➤ 一级和二级缓存存储本地内存、全局内存和寄存器溢出的部分。

内存层次结构

Name	Cache?	Latency (cycle)	Read-only?
全局内存	L1/L2	200~400 (cache miss)	R/W
共享内存	No	1~3	R/W
常量内存	Yes	1~3	Read-only
纹理内存	Yes	~100	Read-only
本地内存	L1/L2	200~400 (cache miss)	R/W



目录

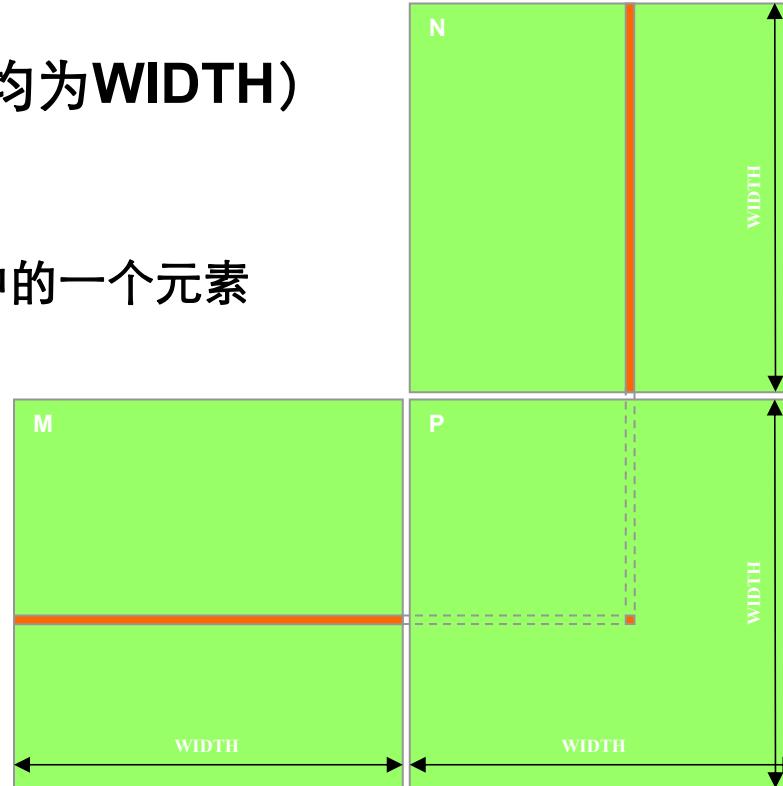
- GPU微处理器简介
- CUDA编程
 - 引言：开发平台、异构计算
 - CUDA编程模型：线程、线程块、线程网格
 - CUDA执行模型：流多处理器和线程束
 - 内存层次结构
 - 编程实例

CUDA程序设计实例——矩阵相乘

- $P = M * N$ (长宽均为WIDTH)

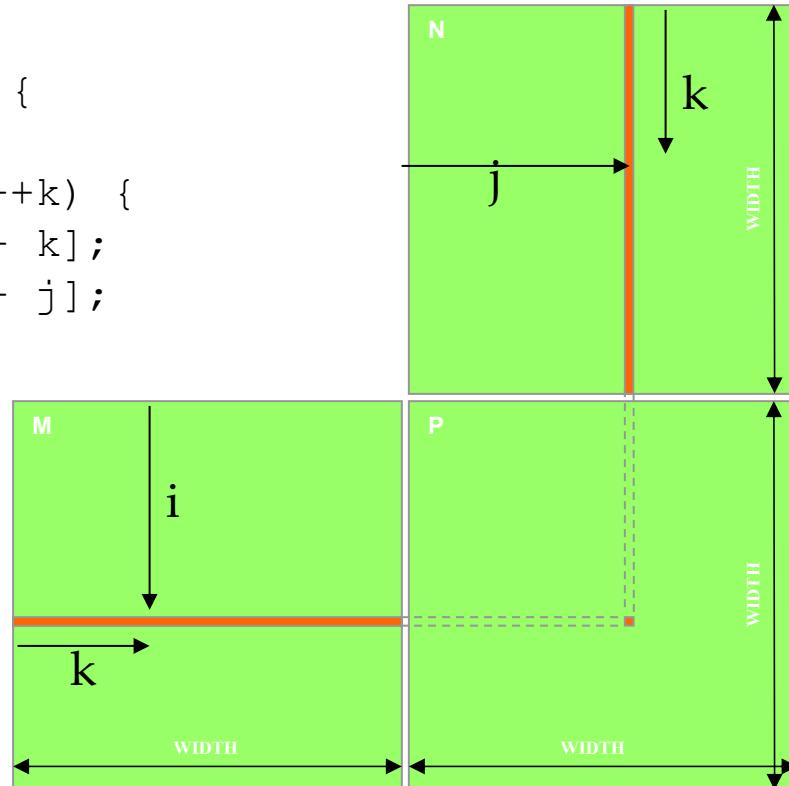
- 算法

每个线程计算矩阵P中的一个元素



第一步：CPU实现

```
// Matrix multiplication on the (CPU) host in double precision
void MatrixMulOnHost(float* M, float* N, float* P, int Width)
{
    for (int i = 0; i < Width; ++i)
        for (int j = 0; j < Width; ++j) {
            double sum = 0;
            for (int k = 0; k < Width; ++k) {
                double a = M[i * width + k];
                double b = N[k * width + j];
                sum += a * b;
            }
            P[i * width + j] = sum;
        }
}
```



第二步：将矩阵数据传给设备（GPU）内存

```
void MatrixMulOnDevice(float* M, float* N, float* P, int
Width)
{
    int size = Width * Width * sizeof(float);
    float* Md, Nd, Pd;
    ...
1. // Allocate and Load M, N to device memory
    cudaMalloc(&Md, size);
    cudaMemcpy(Md, M, size, cudaMemcpyHostToDevice);

    cudaMalloc(&Nd, size);
    cudaMemcpy(Nd, N, size, cudaMemcpyHostToDevice);

    // Allocate P on the device
    cudaMalloc(&Pd, size);
```

第三步：将计算结果传回主机（CPU）内存

```
2. // Kernel invocation code - to be shown later  
...  
  
3. // Read P from the device  
    cudaMemcpy(P, Pd, size,  
cudaMemcpyDeviceToHost);  
  
    // Free device matrices  
    cudaFree(Md) ; cudaFree(Nd) ; cudaFree (Pd) ;  
}
```

第四步： kernel函数

```
// Matrix multiplication kernel - per thread code

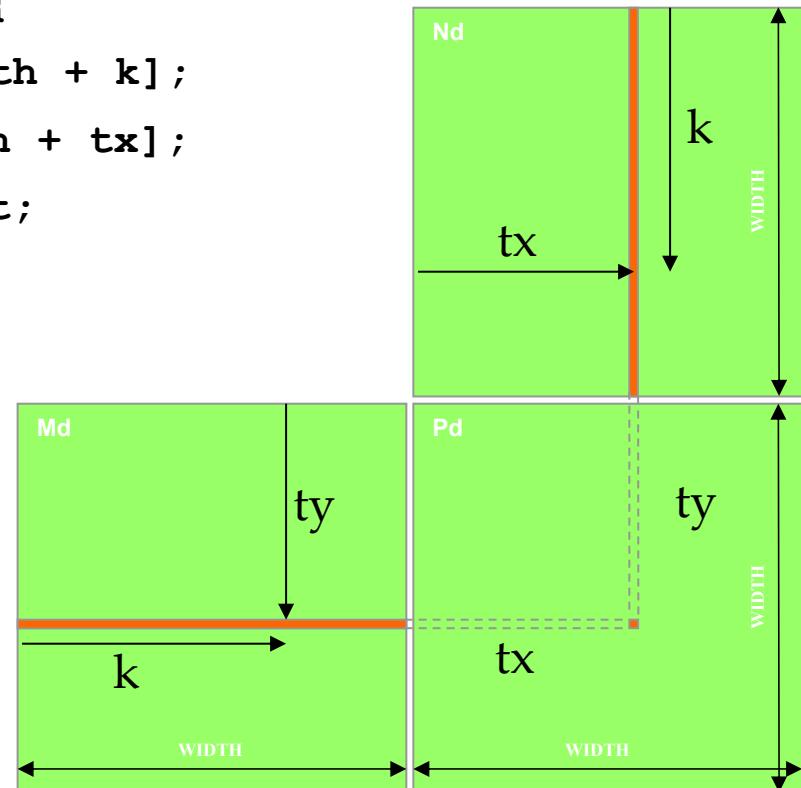
__global__ void MatrixMulKernel(float* Md, float* Nd,
float* Pd, int Width)

{
    // 2D Thread ID
    int tx = threadIdx.x;
    int ty = threadIdx.y;

    // Pvalue is used to store the element of the
matrix
    // that is computed by the thread
    float Pvalue = 0;
```

第四步：kernel函数（续）

```
for (int k = 0; k < Width; ++k)  {  
    float Melement = Md[ty * Width + k];  
    float Nelement = Nd[k * Width + tx];  
    Pvalue += Melement * Nelement;  
}  
  
Pd[ty * Width + tx] = Pvalue;  
}
```



第五步：调用**kernel**函数

```
2. // Kernel invocation code  
    // Setup the execution configuration  
    dim3 dimBlock(Width, Width);  
    dim3 dimGrid(1, 1);  
  
    // Launch the device computation threads!  
    MatrixMulKernel<<<dimGrid, dimBlock>>>(Md, Nd,  
Pd);
```

实例总结：

- 一个线程block计算Pd

- 每个线程计算Pd的一个元素

- 每个线程：

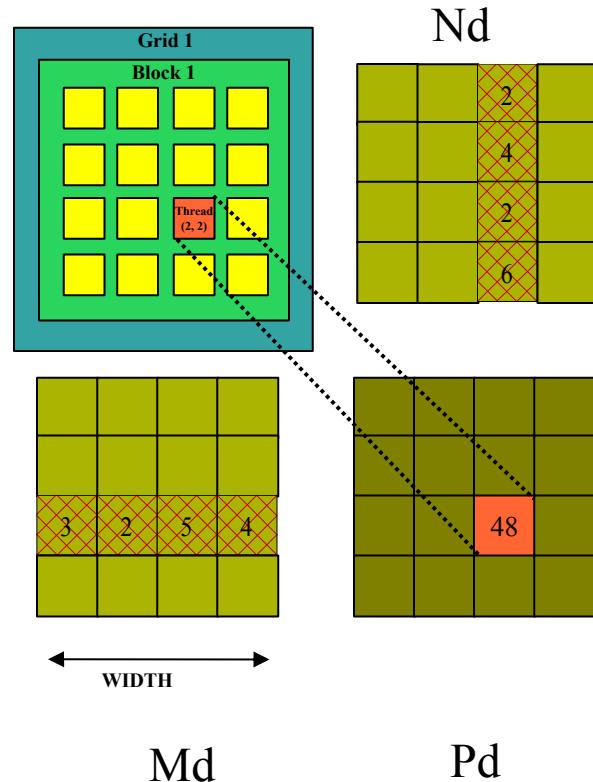
- 读Md矩阵的一行

- 读Nd矩阵的一列

- 为每对Md和Nd元素执行一次乘法和加法；

- 计算次数和片外访存次数比率接近1:1（不是很高）

- 矩阵规模受限于每个block允许的thread数目



课程小结

- GPU微处理器简介
- CUDA编程
 - CUDA编程模型：线程、线程块、线程网格
 - CUDA执行模型：流多处理器和线程束
 - 内存层次结构

推荐网站和读物

- Gerassimos Barlas著，张云泉等译，多核与GPU编程 工具、方法及实践，机械工业出版社，2017
- John Cheng等著，颜成钢等译，CUDA C编程权威指南，机械工业出版社，2017

下一讲

- 大规模数据处理及Map/Reduce编程